



Voxprint AI Audiobook Builder

User manual

Version 0.1.0 (beta) · manual of 3 October 2026

Apache-2.0 · © 2026 Aleksandr Mitroshenkov

Contents


1	What Voxprint does	3
1.1	What you need	4
1.2	How it works in five steps	4
2	Installing Voxprint	5
2.1	The installer	5
2.2	First start.....	5
2.3	Where things are stored	5
3	Quick start	6
4	The Studio (home screen)	7
5	Narrate a book	9
5.1	Step 1 – the book	9
5.2	Step 2 – the voice	10
5.3	Step 3 – prepare the text	10
5.4	Step 4 – output format and quality	11
5.5	Starting, pausing, stopping.....	12
5.6	Speed	13
6	Train your voice	14
6.1	Recording the voice.....	14
6.2	The Train window.....	14
6.3	Training presets	16
6.4	No transcript? (audio only)	17
6.5	Quick preview and comparing two variants.....	17
6.6	The consent card	18
6.7	The universal model (optional)	19
7	My voices.....	20
7.1	Online voices and the voice repository.....	21
7.2	Importing a voice.....	22
7.3	Licences and the two voices that Voxprint offers.....	22
8	Settings.....	23
8.1	Maximum quality (auto).....	24
9	Glossary	26
10	Troubleshooting and FAQ	28
11	Licences and responsibility	30

1 What Voxprint does

Voxprint AI Audiobook Builder ([Train a voice, narrate books](#)) is a Windows 11 application that does two things, entirely on your own computer:

1. **It teaches a computer voice to sound like a particular person.** You give it a recording of a voice (yours, or of a person who has given you permission) – ideally 5 to 15 minutes – and, if you have it, the text that was read. Voxprint aligns text and sound, cuts the recording into clean pieces and trains a small voice add-on (an *adapter*, technically LoRA) for the speech model **Qwen3-TTS**.
2. **It reads whole books aloud with that voice.** You pick a book (TXT, FB2 or EPUB) and a voice from your library, and get a finished audiobook with chapters – one Opus file, one MP3 per chapter, and so on.

There is no command line, no browser and no account. Recordings, texts, voices and results stay on the computer. The program only goes online to download models, to check for updates and, if you ask for it, to download voices from the voice repository.



The Studio – the first screen of Voxprint (dark theme, English interface).

Beta software. Version 0.1.0 has been tested end to end on one machine (Windows with an NVIDIA RTX 4090) and one main speaker. Expect rough edges; keep a copy of your recordings and voices (see [Models and voices](#) in chapter 8).

1.1 What you need

	Minimum	Recommended
Operating system	Windows 11 24H2 (build 26100), 64-bit	Windows 11 26H2
Graphics card	NVIDIA with about 6 GB video memory (VRAM)	NVIDIA with 16 GB VRAM or more
Without an NVIDIA card	The dataset can still be created; training runs on the processor and takes many hours (the program warns you)	–
Disk space	about 12 GB (models about 7 GB + program) plus about 4.2 GB if you build the optional universal model	an SSD
Internet	once, to download the models (about 7 GB)	–

1.2 How it works in five steps

1. Read a known text aloud and record it (5–15 minutes).
2. Choose the recording and the text file – or only the recording, if you have no text.
3. Voxprint aligns the text to the audio with a neural forced aligner and cuts it into clean fragments – the *dataset*.
4. It trains a small voice adapter on your graphics card.
5. The voice appears in [My voices](#). Now [Narrate a book](#) with it.

Use voices responsibly. A voice is personal data and a personality right of its owner. Use a voice only with the explicit permission of the person it belongs to and within the licence of that voice. See chapter 11.

2 Installing Voxprint

2.1 The installer

Run **Voxprint-Setup.exe** (Inno Setup, per-machine installation, about 1.8 GB because the PyTorch libraries are inside). The wizard is available in English, Russian and German. After the install-folder page there is an optional page **Existing models (optional)**: if you already have downloaded models from an earlier installation or from a backup, choose that folder; otherwise leave the field empty. The installer copies nothing – it only remembers the folder, and the first start of the program imports the models from there instead of downloading them.

Silent installation for administrators: `Voxprint-Setup.exe /VERYSILENT /ModelsDir="D:\old\models"`.

If the administrator account that runs the installer differs from the account that uses Voxprint, the data folder belongs to the administrator. In that case set the existing-models folder later in **Settings** (chapter 8).

2.2 First start

On the first start Voxprint prepares itself: the status line says **First start: preparing the program (downloading models, done once)...** It downloads the models it needs (about 7 GB, once). If Hugging Face is slow or unreachable, Voxprint switches to the ModelScope mirror automatically. The download can be interrupted and continues where it stopped. Models already present on the computer (for example from the Hugging Face cache or Alexandria/Pinokio) are reused without downloading them again, read-only.

When it is done, the status shows **Everything is ready. Choose the audio and the text.** A privacy notice is shown on the first start; a short reminder (**Use voices responsibly: only with the voice owner's permission and in line with their rights. Recordings and results stay on this computer.**) stays at the bottom of the windows.

2.3 Where things are stored

What	Where
Program data, models, logs, voice library	%LOCALAPPDATA%\Voxprint\ (models\, logs\, voices\, state\)
Audiobooks	Documents\Voxprint\Audiobooks (can be changed under Advanced in the Narrate window)
Training results	next to your recording: <recording name>_Voxprint\dataset and output\<voice name>

You can open the model folder and the data/log folder with buttons in **Settings** .


3 Quick start

I just want to hear a book (no voice of my own yet). Open [My voices](#) , press [Download voices from repository](#) (or simply wait: online voices are listed by themselves), and download the *Open universal voice* (English, licence CC0 – free for any use). Go back to the Studio, open [Narrate a book](#) , choose a TXT/FB2/EPUB file, pick the voice and press [Start narration](#) .

I want my own voice. Record yourself reading the recording script (section 6.1), open [Train your voice](#) , choose the recording and the script text, leave the training quality at [Balanced \(recommended\)](#) and press [Create voice \(LoRA\)](#) . After training, confirm the consent card. The voice is now in [My voices](#) and in [Narrate a book](#) .

4 The Studio (home screen)

The Studio is the first window. It has no settings of its own – it is a menu of three big cards and a gear.



The Studio.

Element	What it does and when to use it
Title and tagline	"Voxprint AI Audiobook Builder – Train a voice, narrate books ". Only informational.
Narrate a book card	Pick a book (TXT, FB2, EPUB) and a voice - get a finished audiobook with chapters. Opens the narration window. This is the main card. If the library is empty, the card says No voice yet - train your own or import one in "My voices" first. While a book is being narrated it shows Narration in progress...
Train your voice card	Turn a recording of a voice (yours, or of a person who has given permission) and its text into a voice for narration. Opens the training window. While training runs it shows Training in progress...
My voices card	Your voice library: preview, licences, import, download, delete. At the bottom of the card a counter shows how many voices you have (for example "3 voice(s) in your library").




Element	What it does and when to use it
Gear button (top right)	Opens Settings (chapter 8).
Note at the bottom	<i>Use voices responsibly: only with the voice owner's permission and in line with their rights. Recordings and results stay on this computer.</i>

Every other window has a **← Studio** button at the top left that brings you back here. Heavy work (training, narration) keeps running in the background if you go back.

5 Narrate a book

Open this window from the **Narrate a book** card. **Choose a book and a voice, then press Start.** Finished parts are saved on the way, so you can pause, quit and continue later.



The screenshot shows the 'Narrate a book' window. At the top, there's a title bar with a back arrow and 'Studio'. Below it, the title 'Narrate a book' is displayed. A subtitle reads: 'Choose a book and a voice, then press Start. Finished parts are saved on the way, so you can pause, quit and continue later.' The main content area is divided into sections. The first section, '1. Book', contains a 'Choose book...' button and a text field showing 'The Lighthouse Keeper.txt'. Below this, it displays 'The Lighthouse Keeper' and '3 chapters · about 7 min of audio'. The second section is a yellow bar with two options: 'Personal use only · custom/personal-only' and 'Private only'. The third section, '3. Prepare the text', has a subtitle 'Fully automatic - nothing to edit or review. Rules run first, then the optional AI clean-up, then the voice reads the result.' It lists several checkboxes, all of which are checked: 'Tidy the layout' (joins broken lines and hyphenated words, removes odd spaces and invisible characters), 'Remove footnote marks and page numbers' (drops [1], 1, page numbers and headers repeated on every page), 'Normalize quotes and dashes' (unifies quotation marks, dialogue dashes and ellipses), 'Shorten links and e-mail addresses' (a web address is read as "link" instead of being spelled out), 'Fix chapter headings' ('CHAPTER XII' becomes "Chapter twelve"; ALL-CAPS headings become normal text), 'Say numbers in words' (years, dates, percents, money, units and ordinals (Russian and English books)), and 'Expand abbreviations' ('Dr.', 'etc.', 'e.g.' are read in full (Russian and English books)). There is an unchecked checkbox for 'AI: fix typos and add missing commas' with a note 'Small offline model for Russian; every change is checked and only safe ones are applied.' and a note 'Russian books only.' Below this is a link 'More preparation (coming later)'.

The Narrate window, upper half: book, voice and text preparation (demonstration voice).

5.1 Step 1 – the book

Control	What it does
Choose book...	Opens a file dialog. Supported: <code>.txt</code> (UTF-8 or cp1251; headings such as "Chapter 1" or "Глава 2" become chapters), <code>.fb2</code> and <code>.fb2.zip</code> , <code>.epub</code> . Title, author and cover are read when present.
File name and info line	After choosing, Voxprint shows the title and "3 chapters · about 7 min of audio" – the number of chapters and the expected length. Before a choice it says No book chosen (TXT, FB2, EPUB)

If the file cannot be opened you get a message such as **Unsupported book format. Use TXT, FB2 (or FB2.ZIP) or EPUB.** or **The book contains no text.** (see the troubleshooting chapter).

5.2 Step 2 – the voice

The drop-down under **2. Voice** lists the voices of your library. Online voices that are not installed yet appear as " (to download,)"; they are downloaded and verified automatically when you press Start. Below the list a coloured **licence badge** shows what you may do with the result (see chapter 7 and the glossary in chapter 9):

- green – commercial use allowed (for example "CC-BY-4.0 · commercial use OK" and "Commercial OK");
- amber – restricted ("Personal use only · ...", "Private only", "Public, non-commercial").

Under the badge a reminder repeats the limit, for example *Personal use only: do not publish or sell narration made with this voice.* If you have no voices at all, the window says *You have no voices yet. Train your own voice or add one in "My voices".* and offers the buttons **Train your voice** and **My voices**.

5.3 Step 3 – prepare the text

Fully automatic - nothing to edit or review. Rules run first, then the optional AI clean-up, then the voice reads the result. Nothing is shown for review and your book file is never changed. All rule-based steps are ticked by default; untick what you do not want.

Check box	What it does
Tidy the layout	Joins broken lines and hyphenated words, removes odd spaces and invisible characters.
Remove footnote marks and page numbers	Drops [1], ¹ , page numbers and headers repeated on every page.
Normalize quotes and dashes	Unifies quotation marks, dialogue dashes and ellipses.
Shorten links and e-mail addresses	A web address is read as "link" instead of being spelled out.
Fix chapter headings	"CHAPTER XII" becomes "Chapter twelve"; ALL-CAPS headings become normal text.
Say numbers in words	Years, dates, percents, money, units and ordinals (Russian and English books).
Expand abbreviations	"Dr.", "etc.", "e.g." are read in full (Russian and English books).
AI: fix typos and add missing commas	Small offline model for Russian; every change is checked and only safe ones are applied. It needs a one-time download (a button Download appears; the model is about 365 MB) and is available for Russian books only.
► More preparation (coming later)	Greyed-out ideas for later versions: pauses and punctuation by AI, Russian stress marks, translation before narrating, different voices for characters. They are marked " coming later " and do nothing yet.

The rules cover Russian and English fully. For other languages (for example German books) only the language-neutral steps – layout, footnotes, quotes, links – are applied.



5.4 Step 4 – output format and quality

4. Output format

☒ One Opus file with chapters (.opus) — recommended
Open, royalty-free and small. Plays in VLC, Audiobookshelf, Android apps such as Voice.
Free players that handle it: VLC, Audiobookshelf, Voice (Android), Smart AudioBook Player (Android). Chapter support varies by player.

☐ Folder of MP3 files, one per chapter
The most compatible: every player, car, phone and speaker. Includes a playlist and cover.

☐ M4B (AAC) with chapters — for Apple Books
Plays in Apple Books, iPhone/Mac and most audiobook apps.

M4B (AAC) for Apple Books — uses the AAC codec, which is patent-encumbered and not fully open; you choose it at your own discretion.

Quality

Compact
Standard
High

Clean speech and small files - recommended. About Opus 14 MB per hour of audio.

► Other formats

► Advanced

Start narration

Pause Cancel

33%

Fragment 3 of 8 - about 7 min left

Listen to the result so far

Play 0:00 / 0:21

Ready to listen: 3 of 8 parts. New parts are added as they are made.

Narration runs on your computer. Use only voices you are allowed to use.

The Narrate window, lower half: output format, quality, progress and the mini player.

4. Output format – choose one:

Format	What you get	Where it plays
One Opus file with chapters (.opus) — recommended	one file <code>Author - Title.opus</code> with chapter markers	Open, royalty-free and small. Plays in VLC, Audiobookshelf, Android apps such as Voice.
Folder of MP3 files, one per chapter	a folder with <code>NN - Chapter.mp3</code> , a playlist and a cover	The most compatible: every player, car, phone and speaker. Includes a playlist and cover.
M4B (AAC) with chapters — for Apple Books	one <code>.m4b</code> file with chapters	Plays in Apple Books, iPhone/Mac and most audiobook apps.



Choosing M4B shows a yellow patent notice: *M4B (AAC) for Apple Books — uses the AAC codec, which is patent-encumbered and not fully open; you choose it at your own discretion.* In short, AAC is patent-encumbered, Voxprint grants no patent licence for it, and you alone are responsible for legal compliance when you choose it. If you do not need Apple Books, do not use it.

► **Other formats** (collapsed) holds further formats: **M4B container with Opus audio**, **Folder of Opus files, one per chapter**, **One MP3 file with chapter marks**, **Folder of FLAC files (lossless archive)**, **Folder of WAV files (lossless archive)**.

Quality – three buttons that set the bitrate (the amount of data per second of sound – more data, better sound, bigger files):

Button	Meaning	Opus / MP3 / AAC	About per hour of audio
Compact	Smallest files; speech stays clear.	24 / 64 / 48 kbit/s	11 / 29 / 22 MB
Standard	Clean speech and small files - recommended.	32 / 96 / 64 kbit/s	14 / 43 / 29 MB
High	Extra headroom for headphones and noisy places.	48 / 128 / 96 kbit/s	22 / 58 / 43 MB

Pauses – a slider with five positions, from *much shorter* through *normal* (the default) to *much longer*. Voxprint cuts the text at every comma, sentence end, ellipsis, dash, paragraph, scene break and chapter end and inserts **measured silence** there, so the pauses are audible no matter how the voice model phrases the text. At *normal* a sentence is followed by 430 ms, a paragraph by 1.05 s and a chapter by 1.9 s – a little longer than in earlier versions; each kind has its own length (a comma 240 ms, an ellipsis 680 ms, a dash 320 ms, a scene break 1.9 s) and the slider scales them all. The choice is remembered. Changing it for a book that was already narrated re-joins the stored fragments instead of speaking them again, as long as the text pieces stay the same.

► **Advanced** (collapsed) contains: the exact bitrates per format (changing one switches the quality to "custom": no button is highlighted); **Output folder...** – the folder where the audiobook is placed; the check box **Read chapter titles aloud**; and a read-only **Sample of the prepared text (first paragraph)** showing how the first changed paragraph looks after preparation.

5.5 Starting, pausing, stopping

Control	What it does
Start narration	Starts the narration. The first start loads the voice model (<i>Loading the voice model...</i>).
Pause / Resume	Pause stops after the current fragment; the same button then reads Resume.
Cancel	Stops the job. Finished fragments stay saved.
Progress bar and status line	e.g. "Fragment 3 of 8 · about 7 min left" – how many pieces (chunks) of the book are done and the estimated time left.
Open the result folder	Appears at the end (<i>Done</i>): opens the folder with the finished audiobook.

Resumable. Every finished fragment is stored on disk. If you pause, cancel, the program crashes or you close it, press **Start narration** again later: *Continuing: ... fragments were already finished.* The stored fragments are deleted after a successful run.



Listen while it is being made. As soon as the first fragments are ready the **mini player** appears (*Listen to the result so far*): a Play/Pause button, a slider and the time. *Ready to listen: ... of ... parts. New parts are added as they are made.* It follows new fragments and waits if it catches up with the narration. When the job ends, the whole result becomes its playlist.

After a successful run you see *Done.* and, depending on the voice licence, a reminder, for example *Reminder: this voice is private only - keep the audiobook on your computer and do not publish or sell it.*

5.6 Speed

Narration uses **batched generation**: several chunks go through the voice model in one call instead of one at a time. The batch size is chosen from the free video memory (up to 12) and is halved automatically if memory runs out; if a batch fails for any other reason, Voxprint falls back to one chunk at a time, so batching can never make a book fail. Finished chunks are saved by a helper thread while the graphics card already makes the next batch.

Measured on an RTX 4090 (12 English chunks, 80 s of audio, open voice): the old mode needed a real-time factor (RTF) of **3.05** (243 s of work for 80 s of sound); with batch 12 it is **0.31** (25 s) – about **10 times faster**. RTF is "seconds of work per second of audio": 0.31 means one hour of audiobook takes about 19 minutes. Planning figures at RTF 0.31–0.5: 5 hours of audio $\approx$ 1.6–2.5 hours of work; 20 hours $\approx$ 6–10 hours. On cards with 12–16 GB the batch is smaller, so the speed is lower.

6 Train your voice

6.1 Recording the voice

The quality of the voice depends mostly on the recording. Plan **5 to 15 minutes** of clean speech (under 2 minutes triggers a warning; under 10 seconds is refused). The program contains no recorder – use any recorder or phone app and save WAV, FLAC, MP3, M4A or OGG.

Before you record: a quiet room, the microphone about a hand's width from the mouth, no sound "enhancements" (noise suppression, equaliser, reverb), even volume.

The recording script. The repository contains a ready-made script, `docs/voice-script-ru-v3.txt` (Russian; it also has English and German passages). It has 18 blocks: greeting, numbers and dates, sounds of the language, neutral reading, calm, joy, sadness, irritation, surprise, whisper and loud speech, three short stories, optional English, optional technical terms, an optional adults-only block, and – at the very end – the consent statement. The script itself states three simple rules:

1. **Read only ordinary lines.** Everything in square brackets and every line with `===` is a hint – do not read it aloud.
2. **One line – one sentence.** Read it, breathe, stay silent for a second, then read the next line.
3. **Stumbled? Stay silent for a second and read the same sentence again from the beginning.** Do not stop the recording.

Further tips from the script: no need to act the emotions – a slightly different voice is enough, and if it does not work, read in your normal voice. Blocks marked OPTIONAL can be skipped. You choose this same script file as the "text" in the Train window.

Voxprint is built to cope with re-read sentences: when a recording is matched with the script, the best reading of every line is kept. Still, the cleaner the recording, the better the voice.

The consent statement (block 18). Finish the recording with **one** of the consent sentences – the one that is true for you: *commercial* (A), *public non-commercial* (B) or *private only* (C), in Russian, English or German. Say your own name and today's date in place of the bracketed words. Read it only if the voice is yours or you have the owner's permission. Details in section 6.6.

6.2 The Train window

Open it from the `Train your voice` card. **Your voice from a recording in one click. Choose the recording and the text you read.**

Studio
Train your voice

Your voice from a recording in one click. Choose the recording and the text you read.

Choose audio71 audio file(s) chosen

☒ I have no transcript - recognise the speech automatically

Choose a folder with clips...

Automatic speech recognition can make mistakes, and the result is not guaranteed: wrong words in the text teach the voice wrong sounds. Clips the program is unsure about are skipped, but please use clear recordings of one speaker. You can pick many files (or a folder) at once, or drop them here.

☒ I understand that the text will be recognised automatically and may contain errors

Training qualityBalanced (recommended)

The automatic choice for your recording and graphics card. Good likeness in reasonable time.
Estimated training time: about 2 min on NVIDIA GeForce RTX 4090 (5 passes, adapter size 32, about 71 clips).

Advanced

Voice owner's consent

Read it from the spoken statement at the end of my recording (recommended)

End your recording with one of the three consent sentences from the recording script (it says your name, the date and what is allowed). The program recognises it and shows what it understood - you confirm with one click. If it is unclear, the strictest level (private only) is used. This is a record of what you said, not legal advice.

☒ Keep the spoken statement with the voice

Create voice (LoRA)

The small adapter (tens of MB) works only in a few apps, e.g. Alexandria.

Quick preview...

Compare 2 variants

☒ Check the voice after training (recommended)

Quick preview: about 60 s (1 short training run(s) on a few clips, then a ~10 s sample with automatic checks).

Voice type: not set

Optional description of the voice (saved in voice.json)

Build universal model (~4 GB)

The ~4 GB merged model is universal: it works in any app that runs Qwen3-TTS.

Create dataset

0%

UpdatesLoading modelAlignmentSlicingLoRA trainingSaving

Choose the audio and the text - I will do the rest.

The Train window with the "no transcript" mode switched on (English interface).

Control	What it does and when to use it
Choose audio	Choose the voice recording (you can also drag the file into the window). Before a choice: No file chosen (you can drag it into the window)
Choose text	Choose the text file you read (UTF-8 <code>.txt</code> ; the recording script, for example). Before a choice: No file chosen (txt, UTF-8) In no-transcript mode the same row becomes Script (optional): the text that was read aloud... (optional).
Check box I have no transcript - recognise the speech automatically	Audio only, no text: see section 6.4.
Training quality	Drop-down with the training presets: see section 6.3. Below it the program shows a short description and an estimate of the training time for your graphics card.
► Advanced	Shows the manual numbers (only editable in Manual).
Voice owner's consent	How the permission of the voice owner is recorded: see section 6.6.

Control	What it does and when to use it
Create voice (LoRA)	The main button: aligns the recording, cuts it, trains the voice and puts it in your library.
Quick preview... and the check box Compare 2 variants	A quick listen-and-compare test before the long run: section 6.5.
Check box Check the voice after training (recommended)	After training, Voxprint re-reads a test sentence with the new voice and checks pitch and clarity automatically. Leave it on.
Voice type drop-down	Optional: Male voice , Female voice , Child voice , Other voice , or Voice type: not set . Saved with the voice.
Description field	Optional description of the voice (saved in voice.json)
Build universal model (~4 GB)	Merges the trained adapter into one standalone model that works in any app running Qwen3-TTS (section 6.7).
Create dataset	Only prepares the dataset (aligned and cut clips) without training. Useful if you want to train elsewhere.
Progress bar and stage chips	Loading model , Alignment , Slicing , LoRA training , Saving – the stage that is running is highlighted. Cancel appears while a job runs.
Status line	Starts as Choose the audio and the text - I will do the rest.
Gear (top right)	Opens Settings .

When the job ends, the status shows *The voice was added to "My voices"*. and a button **Open the result folder** appears. The adapter folder is `...\output\<voice name>` (tens of MB); the dataset is in `...\<recording name>_Voxprint\dataset`.

In the default mode (with a text) the two buttons **Choose audio** and **Choose text** are shown one under the other, each with the name of the chosen file, and the transcript block with the warning is hidden. (There is no separate screenshot of this default state; the picture above shows the no-transcript mode, which has all the other controls.)

6.3 Training presets

Training quality offers four presets. For each one the program shows what it means and an **estimated time** – for example "about 2 min on NVIDIA GeForce RTX 4090 (5 passes, adapter size 32, about 71 clips)". Read the estimate as "about".

Preset	Meaning
Fast	A quick try: fewer passes over your recording and a smaller adapter. The voice is recognisable, fine details are rougher.
Balanced (recommended)	The automatic choice for your recording and graphics card. Good likeness in reasonable time.
Maximum	The longest run: about twice the passes and a bigger adapter. Best chance of a close voice, but it takes longest and may overfit on short recordings.
Manual	You choose the numbers yourself (see Advanced). Only change them if you know what they do.

Under **Advanced** you see the numbers behind a preset:



Field	Plain meaning
Passes over the recording (epochs)	How many times the training goes over your whole recording. More passes: closer to the voice, but too many can make it "babble".
Adapter size (LoRA rank)	How big the voice add-on is. Larger – more detail, more memory.
Adapter strength (LoRA alpha)	How strongly the add-on influences the base model.
Learning rate	How big each learning step is. The most sensitive number: do not raise it unless you know why.
Steps per update (gradient accumulation)	How many small steps are summed before the model is updated (saves memory).

Editable only in Manual. Higher learning rates and many passes can lower the training number yet make the voice babble and never stop - the automatic value is safest. No preset ever raises the learning rate. In our tests a lower "training number" (loss) did not mean a better voice – always listen to the result.

6.4 No transcript? (audio only)

Tick ***I have no transcript - recognise the speech automatically*** and the program recognises the speech itself (Qwen3-ASR, offline; the recognition model, about 1.9 GB, is downloaded on first use). The text row stays as an **optional script** (***Script (optional): the text that was read aloud...***, or drop a `.txt` into the window): if you choose the text that was read aloud, every recognised piece is matched to it tolerantly (stumbles and re-read lines are handled), so the final text is the correct one, not the recognised one. In addition you have:

- **Choose audio** – now you can pick **many files** at once; the line shows "N audio file(s) chosen";
- **Choose a folder with clips...** – a whole folder of clips, searched with its subfolders;
- drag and drop of files into the window.

Each file is recognised separately; long recordings are cut at pauses into pieces of up to 14 s; everything is merged into one dataset and made equally loud. Skipped automatically: clips under 1.5 s or over 20 s, empty or implausible recognition, duplicates and clips with poor sound. At the end the window reports, for example, "Recognised N file(s): kept X of Y clips = ... min of ... min of audio".

Automatic speech recognition can make mistakes, and the result is not guaranteed: wrong words in the text teach the voice wrong sounds. Clips the program is unsure about are skipped, but please use clear recordings of one speaker. You can pick many files (or a folder) at once, or drop them here.

Because of this, you must tick ***I understand that the text will be recognised automatically and may contain errors*** before the buttons work (the tick resets every time you switch the mode on).

6.5 Quick preview and comparing two variants

Before a long run you can test settings cheaply. **Quick preview...** trains a very short adapter on a few clips (at most 12 clips and 4 passes, about 3 minutes at most) and then makes a sample of about 10 seconds that you can listen to in the player. The estimate is shown next to the button: ***Quick preview: about ... (... short training run(s) on a few clips, then a ~10 s sample with automatic checks).***



The small adapter (tens of MB) works only in a few apps, e.g. Alexandria.

☒ Compare 2 variants

☒ Check the voice after training (recommended)

Quick preview: about 2 min (2 short training run(s) on a few clips, then a ~10 s sample with automatic checks).

Listen

Variant A: 4 passes, adapter size 32 - sample 22 s, recognition errors 15 %, pitch vs. your recording -0.7 semitones - looks good

Variant B: 6 passes, adapter size 64 - sample 32 s, recognition errors 15 %, pitch vs. your recording +1.2 semitones - looks good

Quick preview of two variants, with the player and the automatic metrics.

- With **Compare 2 variants** ticked, two samples are made: A with the current settings and B with a larger adapter and 1.5 times the passes.
- Every variant gets a line of automatic metrics: the length of the sample, "recognition errors" (how many words an automatic listener got wrong when re-reading the sample – lower is better) and "pitch vs. your recording" in semitones (how far the voice is higher or lower than yours – close to 0 is best). A verdict follows: **looks good**, or a warning.
- **Play** plays a variant. **Use these settings** copies its numbers into the **Manual** preset; then press **Create voice (LoRA)** for the full run.
- A preview voice is temporary and is **not** added to **My voices**. The metrics are a hint, not a verdict: always listen.

6.6 The consent card

Every voice carries a record of what its owner allows. Under **Voice owner's consent** choose how it is set:

Option	When to use it
Read it from the spoken statement at the end of my recording (recommended)	You finished the recording with one of the three consent sentences (block 18 of the script). After training the program recognises the statement and shows the result.
I will choose the scope myself	You type the speaker's name (placeholder " Speaker's name ") and choose the scope yourself. Use it only for your own voice or with the owner's permission.
No statement: private use only	No statement. The voice is marked <i>private only</i> .

The check box **Keep the spoken statement with the voice** stores the spoken statement (the last 45 seconds of the recording) next to the voice.

After training in the automatic mode the **consent card** appears:

☒ Keep the spoken statement with the voice

Detected consent - name: Aleksandr Mitroshenkov; date: 2026-10-03; scope: Public, non-commercial: may be published free, not sold.

Recognised statement: "I, Aleksandr Mitroshenkov, today, October third, twenty twenty-six, allow this recording and the characteristics of my voice to be used to train neural networks and to synthesize speech, including free publishing but not selling."

Check it and confirm (or change the scope) with one click.

Public, non-commercial: may be published free, not sold ▼ Confirm

The consent card after training: detected name, date and scope, and the Confirm button.

It shows the *detected* name, date and scope, and the sentence that was recognised. Check it, change the scope in the drop-down if needed and press **Confirm**. If the statement is unclear or missing, the strictest level – **Private only** – is used. This is a record of what was said, **not legal advice**.


Scope in the drop-down	Badge	Meaning	Licence stored
Commercial: the audio may be sold	Commercial OK	The voice owner allowed commercial use: narration made with this voice may be sold.	CC-BY-4.0
Public, non-commercial: may be published free, not sold	Public, non-commercial	Narration made with this voice may be published free of charge, but not sold.	CC-BY-NC-4.0
Private only: results stay on this computer	Private only	Narration made with this voice must stay on your computer; do not publish or sell it.	custom/ personal-only

6.7 The universal model (optional)

The small adapter (tens of MB) works only in a few apps (for example Alexandria). **Build universal model (~4 GB)** merges it into a standalone model folder (`merged_model`) that works in **any** app that runs Qwen3-TTS. Voxprint checks the free disk space and asks first (**Build the universal model**); it is never built automatically. The button is disabled until a voice has been created. Details are in `USAGE.txt` inside the folder.

7 My voices

Open this from the **My voices** card. **Voices are stored on this computer. Each voice has a licence that says what its narration may be used for.**



Old Radio **Personal use only · CC-BY-NC-4.0** **Public, non-commercial**

Alexander **Personal use only · custom/personal-only** **Private only**

Russian · 8 min of speech · 10 epochs · Male voice · by Aleksandr Mitroshenkov

The original Voxprint voice: the first voice ever trained with this software. Test use only.

Test use only: the original Voxprint voice. Do not publish the generated audio and do not use it commercially.

► Preview Narrate with this voice Details... Delete

Vintage radio announcer

Marta **CC-BY-4.0 · commercial use OK** **Commercial OK**

English · 9 min of speech · 12 epochs · Female voice · by Marta

Bright, friendly storyteller

► Preview Narrate with this voice Details... Delete

Old Radio **Personal use only · CC-BY-NC-4.0** **Public, non-commercial**

English · 15 min of speech · 20 epochs · Other voice · by Radio Club

Vintage radio announcer

► Preview Narrate with this voice Details... Delete

The licence covers the voice model only. Use a voice only with its owner's permission and within its licence (for example CC BY-NC = personal use only); you are responsible for lawful use.

My voices: three voice cards (demonstration voices) with licence and scope badges.

Each voice is a card:

Element	Meaning
Name and two badges	The licence badge (green "... · commercial use OK" or amber "Personal use only · ...") and the scope badge (Commercial OK , Public, non-commercial , Private only).
Information line	Language · minutes of speech used for training · epochs (passes) · voice type · author · (for online voices) size.
Description	What the author wrote. A "test use only" voice also shows a warning.
Preview	Plays a short sample of the voice. Stop stops it. If the voice has no sample: <i>This voice has no sample</i>



Element	Meaning
Narrate with this voice	Opens the Narrate window with this voice already chosen.
Details...	Opens <i>Voice details</i> : Name , Author (whose voice it is) , Licence of the voice , Voice type , Description . Press Save .
Delete	Deletes the voice from this computer after a confirmation: <i>Delete the voice "... from this computer? This cannot be undone.</i>

The footer of the window repeats: *The licence covers the voice model only. Use a voice only with its owner's permission and within its licence (for example CC BY-NC = personal use only); you are responsible for lawful use.*

When the library is empty the window says *No voices yet. Train your own voice from a recording, import a voice folder or .zip, or download voices from the repository.* and offers **Train your voice** .

7.1 Online voices and the voice repository

The screenshot displays the Voxprint interface with several voice cards. At the top, there are two yellow cards for 'Old Radio' with licenses 'Personal use only · CC-BY-NC-4.0' and 'Public, non-commercial'. Below these, a card for 'Marta' is shown with a green 'Commercial OK' button and a 'Private only' button. The card details include 'English · 9 min of speech · 12 epochs · Female voice · by Marta · Mitroshenkov' and a description 'Bright, friendly storyteller: the first voice ever trained with this software. Test use only.' Below the description are buttons for 'Preview', 'Narrate with this voice', 'Details...', and 'Delete'. Another card for 'Alexander' is shown with a yellow 'Personal use only · custom/personal-only' button and a 'Private only' button. The card details include 'Russian · by Aleksandr Mitroshenkov · 54 MB' and a description 'The original Voxprint voice: the first voice ever trained with this software. Test use only.' Below the description is a 'Download' button. At the bottom, another 'Old Radio' card is shown with the same licenses as the top ones.

Voices that are available online but not installed yet show a Download button (demonstration data).

- Voices from the online index appear by themselves as cards with the state **Not downloaded yet** , their licence and size and a **Download** button. The download is checked by a SHA-256 checksum and can be resumed.



- **Refresh list** reloads the list. Offline, Voxprint shows the last known list (**No connection - showing the last known list of online voices.**).
- **Download voices from repository** opens the **Voice repository** dialog with **Refresh** , a field for the address of the index (`index.json`) and **Download selected** .

7.2 Importing a voice

Import voice... offers two ways: **From a folder...** (an adapter folder) and **From a .zip file...** (a voice package). Archives are checked for unsafe paths and size limits. A voice that arrives without a declared licence is treated as *personal use only*.

7.3 Licences and the two voices that Voxprint offers

Licence	Commercial use
CC0-1.0, CC-BY-4.0, CC-BY-SA-4.0	allowed (keep attribution / share-alike if the licence asks for it)
CC-BY-NC-4.0, CC-BY-NC-SA-4.0	not allowed
custom/personal-only (default)	not allowed; results stay on your computer
custom/test-use-only	not allowed; for testing only, no public release

- **Open universal voice** (German UI: *Offene Universalstimme*, Russian UI: *Открытый универсальный голос*) – an English voice for people who have no recording of their own. Trained only from the public-domain LJ Speech dataset (reader Linda Johnson, LibriVox); licence **CC0-1.0** – free for any use, also commercial; attribution to the LJ Speech dataset (Keith Ito) is appreciated.
- **Александр (Alexander)** – a male Russian voice, the first one trained with Voxprint. Licence: **personal use only** – no public use of the output, no commercial projects (scope **Private only**).


Both are downloaded separately (not inside the installer) from the voice repository.

The pictures here and in chapter 5 use demonstration voices; in the screenshots the voice "Alexander" still carries an earlier test-licence label. In the released voice list its badge reads "Personal use only" as described here.

The badge is information, not legal advice. The licence covers the voice model only.

8 Settings

Open **Settings** with the gear button (top right of the Studio and the other windows).



Settings.

Control	What it does
Language	The interface language: English, Deutsch, Русский. It changes at once in all windows.
Maximum quality (auto)	Selects every recommended automatic option and downloads the models they need (see below).
Check for updates	Checks whether newer verified versions of components or models exist and installs them (with a safety check and automatic rollback). Voxprint also checks quietly once a week. Messages: <i>Everything is already up to date.</i> , <i>Restart the program to start using the updates.</i>
Open the models folder	Opens the folder with the downloaded models.



Control	What it does
Open the data and log folder	Opens the data and log folder – send the log to the developer if an error repeats.
Models and voices – check box Include my voices	Whether your voice library is part of the backup (on by default).
Back up models and voices...	Copies models (and voices) to a folder or drive you choose. Shows the size and free space first; resumable; identical files are skipped; every copy is checked. A progress bar and a Cancel button appear.
Restore from backup...	Restores from a Voxprint-backup folder. Voices that already exist with different content are never overwritten.
Existing models folder	Models from a previous install or a backup: they are linked (same drive) or copied into Voxprint and checked by checksum before anything is downloaded. Choose folder... picks the folder, Clear forgets it.
Repair the installation	Re-checks Voxprint's own environment and fixes what is broken . Your voices and models are kept.
About	Version, the idea, how it works, open-source components with their licences, a link to the GitHub repository and third-party notices.
Close	Closes the dialog.

If the program notices a damaged installation, a banner offers the repair (**Voxprint's installation needs repair: ...**). In the installed version, if repair does not help, run the Voxprint installer again – voices and models are kept.

When the program finds an older component in an environment it does not own (for example your own Python), it asks first: **Newer component available** with **Upgrade** or **Not now** ; it never changes your environment silently.

8.1 Maximum quality (auto)

Out of the box, everything that can run automatically is **already selected** and marked with a star and the word **recommended** : if you touch nothing, Voxprint works at maximum quality. The **Maximum quality (auto)** button in Settings selects all of these options again (after you have switched some off) and downloads the models they need, with a progress bar. The single check boxes stay in their windows.

- Text preparation (Narrate a book): the seven rule-based steps – layout, footnote marks and page numbers, quotes and dashes, links, chapter headings, numbers in words, abbreviations.
- AI clean-up of typos (Russian books; one-time download of about 365 MB).
- Training window: **Check the voice after training (recommended)** and **Compare 2 variants** (both need the speech-recognition model, about 1.9 GB, downloaded once).
- Always on, without a switch: alignment of the text onto the audio and its plausibility check, the audio quality filter, the recognition filters of the audio-only mode.
- Not included, because they do not exist yet: punctuation model, stress marks, translation, speaker roles.



The training preset stays at **Balanced (recommended)** on purpose: in our tests longer training with a higher learning rate made voices babble, so more is not better there. Instead, after **Compare 2 variants** the better of the two variants is marked **recommended** (verdict first, then fewer recognition errors, then the smaller pitch shift; with no clear difference variant A, which is cheaper). The choice is still yours: always listen.

9 Glossary

Adapter (LoRA)

A small add-on file (tens of MB) that "teaches" the big speech model one particular voice. Technically LoRA – low-rank adaptation. The button **Create voice (LoRA)** creates it.

Universal model

The adapter merged into a full ~4 GB model folder. It works in any app that runs Qwen3-TTS.

Qwen3-TTS

The open speech-synthesis model that Voxprint uses to speak. The base model is Apache-2.0 licensed.

Dataset

The training material Voxprint cuts from your recording: short clips (3–12 s) with their texts, a reference clip and a report.

Alignment

Finding out at what second each word of the text is spoken in the recording. It allows clean cutting at pauses, never in the middle of a word.

Clip / fragment

A short piece of the recording in the dataset. In the narration window "Fragment 3 of 8" means a piece of the book.

Chunk

A piece of the book text of about a sentence or a clause that is sent to the voice model in one go. The UI calls it a "fragment". Finished chunks are saved, which makes narration resumable.

Batch

Several chunks that the voice model makes together in one call. Bigger batches are much faster but use more video memory (up to 12 at a time).

RTF (real-time factor)

Seconds of computing per second of produced audio. 0.31 means one hour of audiobook takes about 19 minutes. Lower is faster.

Epoch / pass

One complete pass of the training over your whole recording (UI: "**Passes over the recording (epochs)**"). The voice list shows how many epochs a voice was trained for.

Preset

A ready-made set of training numbers: **Fast**, **Balanced (recommended)**, **Maximum**, **Manual**.

Rank, alpha, learning rate, gradient accumulation

Technical training numbers shown under **Advanced**; see section 6.3.

Preview

In **My voices**: a short sample of a voice (**Preview**). In the Train window: the *quick preview* – a short test training with a ~10 s sample and automatic checks; nothing is added to your voices.

Recognition errors (WER)

The share of words that an automatic listener (Qwen3-ASR) gets wrong when it transcribes the generated sample. Lower is better.

Pitch, semitone



How high the voice is. A semitone is the smallest step on a piano. "Pitch vs. your recording: +1.2 semitones" means the voice is slightly higher than the original.

Consent

The voice owner's permission, spoken at the end of the recording or entered by hand, saved with the voice. A record, not legal advice.

Scope

What the consent allows: **Commercial OK**, **Public, non-commercial** or **Private only**.

Licence badge / scope badge

Coloured labels on a voice card and in the Narrate window. Green: commercial use allowed. Amber: restricted. The badge is information, not legal advice.

Voice type

An optional label of the voice: male, female, child or other. Saved with the voice.

Mini player

The small player in the Narrate window that plays the finished part of the book while the rest is still being made.

Bitrate (kbit/s)

How much data one second of sound takes. Higher means better quality and bigger files. Speech needs little: Opus 32 kbit/s is plenty.

Opus, MP3, AAC / M4B, FLAC, WAV

Audio formats. Opus: small, open, free. MP3: plays everywhere. AAC in an M4B container: for Apple Books (patent-encumbered). FLAC and WAV: lossless, large files for archiving and editing.

Chapter marks

Bookmarks inside a single audio file that let a player jump from chapter to chapter.

Video memory (VRAM)

The memory of the graphics card. It limits the size of the model and of the batch.

Resume

Continuing a stopped narration or download from where it stopped.

Checksum (SHA-256)

A fingerprint of a file that proves it was downloaded or copied without damage.

Mirror (ModelScope)

An alternative download server that Voxprint uses if Hugging Face is slow or unreachable.

Repair

Re-checks Voxprint's own environment and fixes what is broken.

10 Troubleshooting and FAQ

Which messages can I get, and what do I do?

Message	What to do
The recording is too short (under 10 seconds). At least a few minutes are needed.	Record longer – 5–15 minutes.
Could not read the audio file. Try another format (WAV or MP3).	Convert the recording to WAV or MP3.
The recording and the text do not seem to match: the text is much longer or shorter than the recording. Check that you chose the right files and read the whole text.	Check that you chose the right audio and text and read the whole text.
The recording has long stretches of speech that are not in the text (at the beginning or the end). Trim the extra part or choose another text.	Trim the extra speech from the start or end of the recording, or choose the correct text.
All fragments of the recording turned out too quiet or distorted. Check the microphone and record again.	Check the microphone and level; record again in a quieter place.
Not enough video memory (VRAM).	Close other programs that use the graphics card; Voxprint retries with a lighter load (<i>Not enough video memory - trying again with a lighter load...</i>). As a last resort Continue on the processor (slow) .
Could not download the model {short}. Check your internet connection and try again. You can download the model manually: {url}	Check your internet connection and start again; the download continues where it stopped.
The voice engine failed on fragment {n}. Finished fragments are saved - start again to continue.	Press Start narration again – finished fragments are kept.
ffmpeg is needed for this format but was not found. Use "Repair" in Settings or choose WAV.	Settings → Repair the installation , or choose WAV.
Unsupported book format. Use TXT, FB2 (or FB2.ZIP) or EPUB.	Convert the book to TXT, FB2 or EPUB.
The book file looks unsafe or is too large and was not opened.	The file looks damaged or too large; try another copy.
No NVIDIA graphics card is being used: training will run on the processor and take many hours. A NVIDIA graphics card is needed for a fast result.	Training works without an NVIDIA card but is extremely slow.
The final training loss is low ($\{loss\} < \{threshold\}$): the voice may have turned to "mush" or the speech may not stop. If it sounds bad, record more text and try again.	The voice may babble. Use Balanced (recommended) , record more text, and listen with the quick preview.

The voice sounds higher than mine. A small upward pitch drift (about +2...+3 semitones was seen on a male voice) is a known effect of this kind of training. The automatic check warns above +4. Try the quick preview with both variants and take the one that is closer.

The voice babbles or never stops. Too many passes or too high a learning rate. Use [Balanced \(recommended\)](#), never raise the learning rate by hand.

Can I stop and continue later? Narration: yes, always (section 5.5). Downloads and backups: yes. Training: cancel is possible, but a full run starts over; use the quick preview first.

Can I use my voice for a published audiobook? Only if the consent scope allows it ([Commercial OK](#) to sell; [Public, non-commercial](#) to publish for free). Voices you train without a consent statement are *private only*. You can change the licence under [Details...](#) if you own the voice.

Does Voxprint send my recordings anywhere? No. It sends nothing to the internet except downloading models, checking for updates and, when you open the repository, fetching the voice index and the voices you choose. No telemetry, no account.

Which languages are supported? The interface: English, Deutsch, Русский. Text preparation covers Russian and English fully. The speech model supports further languages, but they are not tested.

Where is the log? [Open the data and log folder](#) in Settings. Send the file from there if an error repeats.

Is the AAC / M4B option safe to use? It is optional and off by default for good reason: AAC is patent-encumbered and you are responsible for compliance (section 5.4).

11 Licences and responsibility

- **Voxprint's own source code is licensed under the Apache License 2.0** (Copyright 2026 Aleksandr Mitroshenkov; see the `LICENSE` and `NOTICE` files).
- **The Apache licence of the code does not cover models and voices.** The models Voxprint downloads (Qwen3-TTS, Qwen3-ASR and the optional text models) keep their authors' licences. Every voice has its own licence, shown as a badge and stored in `voice.json`.
- Third-party components, with their licences, are listed under **About** and in `THIRD_PARTY_NOTICES.md` and the `licenses` folder of the installation. Qt/PySide6 is used under the LGPL-3.0; the bundled ffmpeg is a separate program with its own licence; the AAC codec is patent-encumbered.
- **Use voices responsibly.** A voice is personal data and a personality right. Use a voice only with the explicit permission of its owner and within its licence; no impersonation, fraud or misleading content; label synthetic speech where this is required. You are responsible for the lawful use of the voices and audio you create.

This manual and the licence badges are information, not legal advice. The consent card records what a speaker said; it does not replace the law of your country.

Voxprint AI Audiobook Builder 0.1.0 (beta). Manual of 3 October 2026.

Voxprint AI Audiobook Builder is an independent project and is not affiliated with other products or services with similar names.