



Voxprint

Voxprint AI Audiobook Builder

User manual

Version 0.1.3 (beta) · build 667 “Menuchah” · manual of 9 October 2026

Apache-2.0 · © 2026 Aleksandr Mitroshenkov

Contents

1	What Voxprint does	3
1.1	What you need.....	4
1.2	How it works in five steps	4
2	Installing Voxprint	4
2.1	The installer	4
2.2	First start.....	5
2.3	Where things are stored	5
3	Quick start	6
4	The Studio (home screen)	6
5	Narrate a book	7
5.1	Step 1 – the book	8
5.2	Step 2 – the voice	8
5.3	Step 3 – prepare the text	8
5.4	Step 4 – output format and quality	10
5.5	Starting, pausing, stopping.....	11
5.6	Speed	12
6	Train your voice	12
6.1	Recording the voice	12
6.2	The Train window.....	13
6.3	Training presets	15
6.4	No transcript? (audio only)	16
6.5	Quick preview and comparing two variants.....	16
6.6	The consent card	17
6.7	The universal model (optional)	18
7	My voices.....	18
7.1	Online voices and the voice repository.....	20
7.2	Importing a voice.....	20
7.3	Licences and the voice that Voxprint offers	20
8	Settings.....	21
8.1	Maximum quality out of the box	23
8.2	Check and repair.....	23
9	Glossary	24
10	Troubleshooting and FAQ	26
11	Command line	28
11.1	Where the program is and the rules	28
11.2	Narrate a book.....	28



11.3	Train a voice.....	29
11.4	Back up and restore	29
11.5	Check the installation.....	29


12	Licences and responsibility	29
----	-----------------------------------	----

1 What Voxprint does

Voxprint AI Audiobook Builder (Train a voice, narrate books) is a Windows 11 application that does two things, entirely on your own computer:

1. **It teaches a computer voice to sound like a particular person.** You give it a recording of a voice (yours, or of a person who has given you permission) – ideally 5 to 15 minutes – and, if you have it, the text that was read. Voxprint aligns text and sound, cuts the recording into clean pieces and trains a small voice add-on (an *adapter*, technically LoRA) for the speech model **Qwen3-TTS**.
2. **It reads whole books aloud with that voice.** You pick a book (TXT, FB2 or EPUB) and a voice from your library, and get a finished audiobook with chapters – one Opus file, one MP3 per chapter, and so on.

You need no command line, no browser and no account (a command line exists for scripts and agents – chapter 11). Recordings, texts, voices and results stay on the computer. The program only goes online to download models, to check for updates and, if you ask for it, to download voices from the voice repository.



The Studio – the first screen of Voxprint (build 667).



Beta software. Version 0.1.3 (build 667 “Menuchah”) has run end to end on one machine (Windows 11 with an NVIDIA RTX 4090) with mainly one speaker. Expect rough edges; keep a copy of your recordings and voices (see [Models and voices](#) in chapter 8).

1.1 What you need

	Minimum	Recommended
Operating system	Windows 11 24H2 (build 26100), 64-bit	Windows 11 26H2
Graphics card	NVIDIA with about 6 GB video memory (VRAM)	NVIDIA with 16 GB VRAM or more
Without an NVIDIA card	The dataset can still be created; training runs on the processor and takes many hours (the program warns you)	–
Disk space	about 30 GB for the <i>Full</i> setup (models about 25 GB + components about 3 GB) plus about 4.2 GB if you build the optional universal model	an SSD
Internet	during setup (the libraries) and once for the models (about 28 GB together with <i>Full</i> ; 1.7B speech recognition on GPUs with 8 GB of video memory or more)	–

1.2 How it works in five steps

1. Read a known text aloud and record it (5–15 minutes).
2. Choose the recording and the text file – or only the recording, if you have no text.
3. Voxprint aligns the text to the audio with a neural forced aligner and cuts it into clean fragments – the *dataset*.
4. It trains a small voice adapter on your graphics card.
5. The voice appears in [My voices](#). Now [Narrate a book](#) with it.

Use voices responsibly. A voice is personal data and a personality right of its owner. Use a voice only with the explicit permission of the person it belongs to and within the licence of that voice. See chapter 12.

2 Installing Voxprint

2.1 The installer

Run **Voxprint-Setup-online.exe** (Inno Setup, per-machine installation, about 34 MB; build 667 has 35,235,971 bytes and the SHA-256 `c145a9fc84c736d655fbb9bfd5c1cc94794b1be7fb8e9a5eec293e6f1578e42`). It is an *online* installer: the libraries (PyTorch and the rest) are downloaded from PyTorch/PyPI during setup. The installer is not code-signed, so Windows SmartScreen warns – choose *More info* → *Run anyway* only after the SHA-256 matches the one on the release page. The wizard is available in English, Russian and German; the language you pick there becomes the program's interface language.



The page **Setup type** offers two choices. **Full** (default) downloads everything on the first start without another click – the components and all models, about 28 GB; the wizard shows the total and checks the free space first. **Quick** installs only the program; the first start opens the **Components** window, which offers the same complete download and starts it when you press Download. Nothing is left out in either mode.

After the install-folder page there is an optional page **Existing models (optional)**: if you already have downloaded models from an earlier installation or from a backup, choose that folder; otherwise leave the field empty. The installer copies nothing – it only remembers the folder, and the first start of the program imports the models from there instead of downloading them.

Silent installation for administrators: `Voxprint-Setup-online.exe /VERYSILENT /Mode=full /ModelsDir="D:\old\models"` (`/Mode=quick` for the Quick type).

If the administrator account that runs the installer differs from the account that uses Voxprint, the data folder belongs to the administrator. In that case set the existing-models folder later in **Settings** (chapter 8).

2.2 First start

On the first start Voxprint prepares itself. With the online installer the **Components** window opens and works without a click: one line shows what is being downloaded right now (**Downloading: {name} — {pct} % · {speed}**), then the checksum check, and a bar shows the overall progress. Step 1 brings the program components together with the SAGE text clean-up model and the translation models, step 2 right after it ALL models (voice, alignment, speech recognition) into your models folder. Otherwise the status line says **First start: preparing the program (downloading models, done once)...** and downloads the models it needs (up to about 15 GB, once). If Hugging Face is slow or unreachable, Voxprint switches to the ModelScope mirror automatically. The download can be interrupted and continues where it stopped. Models already present on the computer (for example from the Hugging Face cache or Alexandria/ Pinokio) are reused without downloading them again, read-only. The two bundled voices *Boaz* and *Tirzah* (chapter 7) are part of this download.

When it is done, the status shows **Everything is ready. Choose the audio and the text.** A privacy notice is shown on the first start; a short reminder (**Use voices responsibly: only with the voice owner's permission and in line with their rights. Recordings and results stay on this computer.**) stays at the bottom of the windows.

2.3 Where things are stored

What	Where
Program data, models, logs, voice library	%LOCALAPPDATA%\Voxprint\ (models\, logs\, voices\, state\)
Projects: audiobooks, voice trainings, Re-voice	the projects folder %LOCALAPPDATA%\Voxprint\Projects with Audiobooks\<book>, Voices\<voice name>_Voxprint and Re-voice; a <i>Voxprint Projects</i> shortcut is in Documents. Change it under Projects folder in Settings .
Training results	Projects\Voices\<voice name>_Voxprint: the dataset and the adapter in output\<voice name>

You can open the model folder and the data/log folder with buttons in **Settings**.


3 Quick start

I just want to hear a book (no voice of my own yet). Open **My voices**, press **Download voices from repository** (or simply wait: online voices are listed by themselves), and download the *Open universal voice* (English, licence CC0 – free for any use). Go back to the Studio, open **Narrate a book**, choose a TXT/FB2/EPUB file, pick the voice and press **Start narration**.

I want my own voice. Record yourself reading the recording script (section 6.1), open **Train your voice**, choose the recording and the script text, leave the training quality at **Balanced (recommended)** and press **Create voice (LoRA)**. After training, confirm the consent card. The voice is now in **My voices** and in **Narrate a book**.

4 The Studio (home screen)

The Studio is the first window. It has no settings of its own – it is a menu of four big cards and a gear.



The Studio.

Element	What it does and when to use it
Title and tagline	"Voxprint AI Audiobook Builder – Train a voice, narrate books ". Only informational.



Element	What it does and when to use it
Narrate a book card	Pick a book (TXT, FB2, EPUB) and a voice - get a finished audiobook with chapters. Opens the narration window. This is the main card. If the library is empty, the card says No voice yet - train your own or import one in "My voices" first. While a book is being narrated it shows Narration in progress...
Train your voice card	Turn a recording of a voice (yours, or of a person who has given permission) and its text into a voice for narration. Opens the training window. While training runs it shows Training in progress...
My voices card	Your voice library: preview, licences, import, download, delete. At the bottom of the card a counter shows how many voices you have (for example "2 voice(s) in your library" – the two bundled voices).
Re-voice card	Record or choose audio, then turn it into editable text and voice it, or re-voice the recording directly in one of your voices.
Gear button (top right)	Opens Settings (chapter 8).
Note at the bottom	Use voices responsibly: only with the voice owner's permission and in line with their rights. Recordings and results stay on this computer.

Every other window has a **← Studio** button at the top left that brings you back here. Heavy work (training, narration) keeps running in the background if you go back.

5 Narrate a book

Open this window from the **Narrate a book** card. **Choose a book and a voice, then press Start.** Finished parts are saved on the way, so you can pause, quit and continue later.

← Studio

Narrate a book

Choose a book and a voice, then press Start. Finished parts are saved on the way, so you can pause, quit and continue later.

1. Book

Choose book... genesis-ru.txt

genesis-ru
2 chapters · about 4 min of audio

2. Voice

Boaz

CC0-1.0 · commercial use OK Commercial OK

Russian · Mark Chulsky (LibriVox) · public domain

☐ Say at the start that the book is narrated by AI (voice name, month and year)

3. Prepare the text

Fully automatic - nothing to edit or review. Rules run first, then the optional AI clean-up, then the voice reads the result.

☒ Tidy the layout ★ recommended
Joins broken lines and hyphenated words, removes odd spaces and invisible characters.

☒ Remove footnote marks and page numbers ★ recommended
Drops [1], ¹, page numbers and headers repeated on every page.

☒ Normalize quotes and dashes ★ recommended
Unifies quotation marks, dialogue dashes and ellipses.

☒ Shorten links and e-mail addresses ★ recommended
A web address is read as "link" instead of being spelled out.

☒ Fix chapter headings ★ recommended
"CHAPTER XIII" becomes "Chapter twelve"; ALL-CAPS headings become normal text.

☒ Say numbers in words ★ recommended
Years, dates, percents, money, units and ordinals (Russian and English books).

☒ Expand abbreviations ★ recommended

4. Output format

☒ One Opus file with chapters (opus) — recommended
Open, royalty-free and small. Plays in VLC, Audiobookshelf, Android apps such as Voice.
Free players that handle it: VLC, Audiobookshelf, Voice (Android), Smart Audiobook Player (Android). Chapter support varies by player.

☐ Folder of MP3 files, one per chapter
The most compatible: every player, car, phone and speaker. Includes a playlist and cover.

☐ M4B (AAC) with chapters — for Apple Books
Plays in Apple Books, iPhone/Mac and most audiobook apps.

☐ M4B (AAC) for Apple Books — uses the AAC codec, which is patent-encumbered and not fully open; you choose it at your own discretion.

Quality

Compact Standard High

Translate the book (optional)

☐ Translate the book before narrating, to: English

Automatic machine translation: quality varies - names, idioms, dialogue and long sentences can come out wrong, and poems lose their rhyme. It runs offline on this computer (graphics card if there is one, otherwise the processor; a long book can take a while). The translated text is saved next to the audiobook as translation_xx.txt: read it, edit it and start again to narrate your edited text. You are responsible for having the right to translate and narrate the text (copyright and licences); this is not legal advice.

AI text model (optional)

Gemma 4 12B is downloaded. It runs before the narration and is unloaded before the voice starts.

☐ Literary translation (AI model, slower; Opus-MT takes over where it fails)

☒ Prepare text for narration (numbers, dates and abbreviations as words, no brackets)

Both are off by default. The model can be wrong: the translation file next to the audiobook can be edited.

Edit prompts

The Narrate window, upper half: book, voice (the bundled voice Boaz), text preparation, translation and the optional AI text model.



5.1 Step 1 – the book

Control	What it does
Choose book...	Opens a file dialog. Supported: <code>.txt</code> (UTF-8 or cp1251; headings such as "Chapter 1" or "Глава 2" become chapters), <code>.fb2</code> and <code>.fb2.zip</code> , <code>.epub</code> . Title, author and cover are read when present.
File name and info line	After choosing, Voxprint shows the title and "3 chapters · about 7 min of audio" – the number of chapters and the expected length. Before a choice it says No book chosen (TXT, FB2, EPUB)

If the file cannot be opened you get a message such as **Unsupported book format. Use TXT, FB2 (or FB2.ZIP) or EPUB.** or **The book contains no text.** (see the troubleshooting chapter).

5.2 Step 2 – the voice

The drop-down under **2. Voice** lists the voices of your library. Online voices that are not installed yet appear as " (to download,)"; they are downloaded and verified automatically when you press Start. Below the list a coloured **licence badge** shows what you may do with the result (see chapter 7 and the glossary in chapter 9):

- green – commercial use allowed (for example "CC0-1.0 · commercial use OK" and "Commercial OK" for the bundled voices);
- amber – restricted ("Personal use only · ...", "Private only", "Public, non-commercial").

Under the badge a reminder repeats the limit, for example **Personal use only: do not publish or sell narration made with this voice.** If you have no voices at all, the window says **You have no voices yet. Train your own voice or add one in "My voices".** and offers the buttons **Train your voice** and **My voices**.

5.3 Step 3 – prepare the text

Fully automatic - nothing to edit or review. Rules run first, then the optional AI clean-up, then the voice reads the result. Nothing is shown for review and your book file is never changed. All rule-based steps are ticked by default; untick what you do not want.

Check box	What it does
Tidy the layout	Joins broken lines and hyphenated words, removes odd spaces and invisible characters.
Remove footnote marks and page numbers	Drops [1], ¹, page numbers and headers repeated on every page.
Normalize quotes and dashes	Unifies quotation marks, dialogue dashes and ellipses.
Shorten links and e-mail addresses	A web address is read as "link" instead of being spelled out.
Fix chapter headings	"CHAPTER XII" becomes "Chapter twelve"; ALL-CAPS headings become normal text.
Say numbers in words	Years, dates, percents, money, units and ordinals (Russian and English books).
Expand abbreviations	"Dr.", "etc.", "e.g." are read in full (Russian and English books).



Check box	What it does
AI: fix typos and add missing commas	Small offline model for Russian; every change is checked and only safe ones are applied. It needs a one-time download (a button Download appears; the model is about 365 MB) and is available for Russian books only.
► More preparation (coming later)	Greyed-out ideas for later versions: pauses and punctuation by AI, Russian stress marks, different voices for characters. They are marked " coming later " and do nothing yet.

The rules cover Russian and English fully. For other languages (for example German books) only the language-neutral steps – layout, footnotes, quotes, links – are applied.

Optional – translate the book

The card **Translate the book (optional)** has a check box (*Translate the book before narrating, to:*) and a list of target languages. Voxprint detects the language of the book itself and translates it sentence by sentence **offline on your computer** with the open Opus-MT models (English ↔ Russian, English ↔ German; Russian ↔ German goes through English and is slower and less accurate). The models are about 300 MB per direction and are downloaded once, with a SHA-256 check, when you press **Download**. A graphics card is used if there is one, otherwise the processor. Chapters, chapter titles, paragraphs, verse lines, scene breaks and therefore the pauses stay as they are; then the usual text preparation runs for the new language and the book is narrated in it – choose a voice of the target language for the best result.

- The translated text is saved next to the audiobook as `translation_xx.txt` (in the folder `Title (xx)`). Open it, read it and edit it; start the same narration again and **your edited text** is narrated. Delete the file to translate anew. Translated sentences are cached, so a stopped job never translates the same sentence twice.
- *Automatic machine translation: quality varies - names, idioms, dialogue and long sentences can come out wrong, and poems lose their rhyme. It runs offline on this computer (graphics card if there is one, otherwise the processor; a long book can take a while). The translated text is saved next to the audiobook as translation_xx.txt: read it, edit it and start again to narrate your edited text. You are responsible for having the right to translate and narrate the text (copyright and licences); this is not legal advice.*
- Machine translation is not a human translation: expect wrong names, idioms, tenses and gender; short titles suffer most. Poems lose their rhyme.

Optional – the AI text model

The card **AI text model (optional)** belongs to the optional Gemma 4 12B text model (about 7 GB with the llama.cpp runtime; it is fetched by the *Full* setup or with `voxprint models download llm`). When it is installed the card says **Gemma 4 12B is downloaded. It runs before the narration and is unloaded before the voice starts.** It offers two check boxes: **Literary translation (AI model, slower; Opus-MT takes over where it fails)** and **Prepare text for narration (numbers, dates and abbreviations as words, no brackets)**. **Both are off by default. The model can be wrong: the translation file next to the audiobook can be edited.** **Edit prompts** opens the folder with the instructions for the model; edited files are used instead of the built-in ones.

Known issue in build 667: the two check boxes of this card may not react to a click even though the model is downloaded. The rule-based text preparation (step 3) is not affected. This will be fixed in a coming build.



5.4 Step 4 – output format and quality

3. Prepare the text

Fully automatic - nothing to edit or review. Rules run first, then the optional AI clean-up, then the voice reads the result.

- ☒ **Tidy the layout** ★ recommended
Joins broken lines and hyphenated words, removes odd spaces and invisible characters.
- ☒ **Remove footnote marks and page numbers** ★ recommended
Drops [1], ¹, page numbers and headers repeated on every page.
- ☒ **Normalize quotes and dashes** ★ recommended
Unifies quotation marks, dialogue dashes and ellipses.
- ☒ **Shorten links and e-mail addresses** ★ recommended
A web address is read as "link" instead of being spelled out.
- ☒ **Fix chapter headings** ★ recommended
"CHAPTER XII" becomes "Chapter twelve"; ALL-CAPS headings become normal text.
- ☒ **Say numbers in words** ★ recommended
Years, dates, percents, money, units and ordinals (Russian and English books).
- ☒ **Expand abbreviations** ★ recommended
"Dr.", "etc.", "e.g." are read in full (Russian and English books).
- ☒ **AI: fix typos and add missing commas** ★ recommended
Small offline model for Russian; every change is checked and only safe ones are applied.
Model downloaded

► More preparation (coming later)

4. Output format

- ☒ **One Opus file with chapters (.opus)** — recommended
Open, royalty-free and small. Plays in VLC, Audiobookshelf, Android apps such as Voice.
Free players that handle it: VLC, Audiobookshelf, Voice (Android), Smart Audiobook Player (Android). Chapter support varies by player.
- ☐ **Folder of MP3 files, one per chapter**
The most compatible: every player, car, phone and speaker. Includes a playlist and cover.
- ☐ **M4B (AAC) with chapters** — for Apple Books
Plays in Apple Books, iPhone/Mac and most audiobook apps.
M4B (AAC) for Apple Books — uses the AAC codec, which is patent-encumbered and not fully open; you choose it at your own discretion.

Quality

Compact Standard High

Clean speech and small files - recommended. About Opus 14 MB per hour of audio.

☐ Check every fragment by speech recognition and regenerate the ones that do not say their text (slower)

► Other formats

☐ Insert measured pauses after commas and sentences (may make the voice swallow short words)

Pauses normal (sentence 430 ms, paragraph 1050 ms)

Silence of a measured length is inserted after every comma, sentence, ellipsis, dash, paragraph, scene break and chapter, independent of how the voice model phrases the text, so the pauses are clearly audible. "Normal" is a little longer than the earlier narration. Changing it later only re-joins the audio when the text pieces stay the same.

► Advanced

Start narration

Choose a book and a voice.

Narration runs on your computer. Use only voices you are allowed to use.

The Narrate window, lower half: text preparation, output format, quality, the measured-pauses option and the Start button.

4. Output format – choose one:

Format	What you get	Where it plays
One Opus file with chapters (.opus) — recommended	one file <code>Author - Title.opus</code> with chapter markers	Open, royalty-free and small. Plays in VLC, Audiobookshelf, Android apps such as Voice.
Folder of MP3 files, one per chapter	a folder with <code>NN - Chapter.mp3</code> , a playlist and a cover	The most compatible: every player, car, phone and speaker. Includes a playlist and cover.
M4B (AAC) with chapters — for Apple Books	one <code>.m4b</code> file with chapters	Plays in Apple Books, iPhone/Mac and most audiobook apps.

Choosing M4B shows a yellow patent notice: **M4B (AAC) for Apple Books — uses the AAC codec, which is patent-encumbered and not fully open; you choose it at your own discretion.** In short, AAC is patent-encumbered, Voxprint grants no patent licence for it, and you alone are responsible for legal compliance when you choose it. If you do not need Apple Books, do not use it.

► **Other formats** (collapsed) holds further formats: **M4B container with Opus audio**, **Folder of Opus files, one per chapter**, **One MP3 file with chapter marks**, **Folder of FLAC files (lossless archive)**, **Folder of WAV files (lossless archive)**.



Quality – three buttons that set the bitrate (the amount of data per second of sound – more data, better sound, bigger files):

Button	Meaning	Opus / MP3 / AAC	About per hour of audio
Compact	Smallest files; speech stays clear.	24 / 64 / 48 kbit/s	11 / 29 / 22 MB
Standard	Clean speech and small files - recommended.	32 / 96 / 64 kbit/s	14 / 43 / 29 MB
High	Extra headroom for headphones and noisy places.	48 / 128 / 96 kbit/s	22 / 58 / 43 MB

Insert measured pauses after commas and sentences (may make the voice swallow short words) – off by default. Ticking it cuts the text at **every** comma as well and inserts measured silence there; the pauses are the clearest, but the voice may swallow short words. Only then is the slider **Pauses** active: five positions from **much shorter** through **normal** (the default) to **much longer**; at **normal** a sentence is followed by 430 ms and a paragraph by 1050 ms, and the slider scales every kind of pause. Changing it for a book that was already narrated re-joins the stored fragments instead of speaking them again, as long as the text pieces stay the same.

Check every fragment by speech recognition and regenerate the ones that do not say their text (slower) – off by default. After synthesis every fragment is re-read by speech recognition; fragments that skip or repeat words are generated again.

Pauses and reading speed

Since version 0.1.3 every narration follows the punctuation and the structure of the text, even with the box above unticked. Voxprint cuts the text per sentence and at strong transitions inside a sentence (dash, colon, semicolon, a comma before "and" / "but"), trims every spoken piece of its own silence and inserts measured silence: comma 0.25 s, strong break 0.40 s, sentence 0.60 s, paragraph or verse line 1.00 s, chapter title, chapter end or scene break 2.00 s. Plain commas stay inside a piece and no piece is shorter than 20 characters, so the voice does not swallow words.

The **reading speed adapts to the text**: long, comma-rich and descriptive sentences and scripture-like text are read a little slower; dialogue and short lines at the voice's own speed. The five pause lengths, a global **Reading speed** (70–130 %) and the **Reading style** (**Automatic**, **Solemn / scripture**, **Fiction**, **Dialogue-heavy**) are set in **Settings** → **Narration: pauses and speed** (chapter 8) and apply to the next narration. The speed is applied after synthesis by time-stretching with the pitch kept, so changing pauses or speed never synthesizes finished fragments again – the next run only joins them anew.

► **Advanced** (collapsed) contains: the exact bitrates per format (changing one switches the quality to "custom": no button is highlighted); **Working folder...** – the folder where the audiobook is placed; the check box **Read chapter titles aloud**; and a read-only **Sample of the prepared text (first paragraph)** showing how the first changed paragraph looks after preparation.

5.5 Starting, pausing, stopping

Control	What it does
Start narration	Starts the narration. The first start loads the voice model (<i>Loading the voice model...</i>).
Pause / Resume	Pause stops after the current fragment; the same button then reads Resume.
Cancel	Stops the job. Finished fragments stay saved.



Control	What it does
Progress bar and status line	e.g. "Fragment 3 of 8 · about 7 min left" – how many pieces (chunks) of the book are done and the estimated time left.
Open the result folder	Appears at the end (Done): opens the folder with the finished audiobook.

Resumable. Every finished fragment is stored on disk. If you pause, cancel, the program crashes or you close it, press **Start narration** again later: **Continuing: ... fragments were already finished.** The stored fragments are deleted after a successful run.

Listen while it is being made. As soon as the first fragments are ready the **mini player** appears (**Listen to the result so far**): a Play/Pause button, a slider and the time. **Ready to listen: ... of ... parts. New parts are added as they are made.** It follows new fragments and waits if it catches up with the narration. When the job ends, the whole result becomes its playlist.

After a successful run you see **Done.** and, depending on the voice licence, a reminder, for example **Reminder: this voice is private only - keep the audiobook on your computer and do not publish or sell it.**

5.6 Speed

Narration uses **batched generation**: several chunks go through the voice model in one call instead of one at a time. The batch size is chosen from the free video memory (up to 12) and is halved automatically if memory runs out; if a batch fails for any other reason, Voxprint falls back to one chunk at a time, so batching can never make a book fail. Finished chunks are saved by a helper thread while the graphics card already makes the next batch.

Measured on an RTX 4090 (12 English chunks, 80 s of audio, open voice): the old mode needed a real-time factor (RTF) of **3.05** (243 s of work for 80 s of sound); with batch 12 it is **0.31** (25 s) – about **10 times faster**. RTF is "seconds of work per second of audio": 0.31 means one hour of audiobook takes about 19 minutes. Planning figures at RTF 0.31–0.5: 5 hours of audio $\approx$ 1.6–2.5 hours of work; 20 hours $\approx$ 6–10 hours. On cards with 12–16 GB the batch is smaller, so the speed is lower.

6 Train your voice

6.1 Recording the voice

The quality of the voice depends mostly on the recording. Plan **5 to 15 minutes** of clean speech (under 2 minutes triggers a warning; under 10 seconds is refused). The program contains no recorder – use any recorder or phone app and save WAV, FLAC, MP3, M4A or OGG.

Before you record: a quiet room, the microphone about a hand's width from the mouth, no sound "enhancements" (noise suppression, equaliser, reverb), even volume.

The recording script. The repository contains ready-made scripts in `docs/recording-scripts/`, written separately in Russian, English and German (TXT and PDF, dark and print). Each has 17 blocks: greeting, the story of the idea, numbers and dates, smooth sentences for the sounds of the language, neutral reading, calm, joy, sadness, irritation, surprise, whisper and loud speech, three short computer stories,



an optional "bad day" story with mild colour words (no profanity), optional technical terms, and – at the very end – the consent statement in the language of the file. The PDFs add a reading guide (tempo, pauses, emotions, breaks). The script itself states three simple rules:

1. **Read only ordinary lines.** Everything in square brackets and every line with `===` is a hint – do not read it aloud.
2. **One line – one sentence.** Read it, breathe, stay silent for a second, then read the next line.
3. **Stumbled? Stay silent for a second and read the same sentence again from the beginning.** Do not stop the recording.

Further tips from the script: no need to act the emotions – a slightly different voice is enough, and if it does not work, read in your normal voice. Blocks marked OPTIONAL can be skipped. You choose this same script file as the "text" in the Train window.

Voxprint is built to cope with re-read sentences: when a recording is matched with the script, the best reading of every line is kept. Still, the cleaner the recording, the better the voice.

The consent statement (block 18). Finish the recording with **one** of the consent sentences – the one that is true for you: *commercial* (A), *public non-commercial* (B) or *private only* (C), in Russian, English or German. Say your own name and today's date in place of the bracketed words. Read it only if the voice is yours or you have the owner's permission. Details in section 6.6.

6.2 The Train window

Open it from the **Train your voice** card. **Your voice from a recording in one click. Choose the recording and the text you read.**

The Train window, upper part: audio and text, training quality, the voice owner's consent and the main button.



Control	What it does and when to use it
Choose audio	Choose the voice recording (you can also drag the file into the window). Before a choice: No file chosen (you can drag it into the window)
Choose text	Choose the text file you read (UTF-8 .txt ; the recording script, for example). Before a choice: No file chosen (txt, UTF-8) In no-transcript mode the same row becomes Script (optional): the text that was read aloud... (optional).
Check box I have no transcript - recognise the speech automatically	Audio only, no text: see section 6.4.
Training quality	Drop-down with the training presets: see section 6.3. Below it the program shows a short description and an estimate of the training time for your graphics card.
► Advanced	Shows the manual numbers (only editable in Manual).
Voice owner's consent	How the permission of the voice owner is recorded: see section 6.6.
Create voice (LoRA)	The main button: aligns the recording, cuts it, trains the voice and puts it in your library.
Quick preview... and the check box Compare 2 variants	A quick listen-and-compare test before the long run: section 6.5.
► Checks and clip options – check box Check the voice after training (recommended)	After training, Voxprint re-reads a test sentence with the new voice and checks pitch and clarity automatically. Leave it on.
Gender: not specified and Age: not specified , ► More about the voice (optional)	Optional details of the voice. The gender (Male voice , Female voice , Child voice , Other voice) is saved in voice.json and added to the name of the trained voice's folder (<code>anna_male</code> , <code>anna_female</code> , <code>anna_unspecified</code>). The program suggests male or female from the pitch of your recording; change it if that is wrong. More about the voice (optional) holds further optional fields for voice.json.
Description field	Optional description of the voice (saved in voice.json)
Build universal model (~4 GB)	Merges the trained adapter into one standalone model that works in any app running Qwen3-TTS (section 6.7).
Create dataset	Only prepares the dataset (aligned and cut clips) without training. Useful if you want to train elsewhere.
Progress bar and stage chips	Loading model , Alignment , Slicing , LoRA training , Saving – the stage that is running is highlighted. Cancel appears while a job runs.
Status line	Starts as Choose the audio and the text - I will do the rest.
Gear (top right)	Opens Settings .



► Advanced

Voice owner's consent

Read it from the spoken statement at the end of my recording (recommended) ▼

End your recording with one of the three consent sentences from the recording script (it says your name, the date and what is allowed). The program recognises it and shows what it understood - you confirm with one click. If it is unclear, the strictest level (private only) is used. This is a record of what you said, not legal advice.

☒ Keep the spoken statement with the voice

Create voice (LoRA)

The small adapter (tens of MB) works only in a few apps, e.g. Alexandria.

Quick preview... ☒ Compare 2 variants ★ recommended

► Checks and clip options

Quick preview: about 3 min (2 short training run(s) on a few clips, then a ~10 s sample with automatic checks).

Voice name (e.g. Anna); empty = from the recording Optional description of the voice (saved in voice.json)

Gender: not specified ▼ Age: not specified ▼ ► More about the voice (optional)

The voice type goes into voice.json and into the name of the trained voice's folder (for example anna_male).

Build universal model (~4 GB)

The ~4 GB merged model is universal: it works in any app that runs Qwen3-TTS.

Create dataset

0%

Updates Loading model Alignment Slicing LoRA training Saving

Choose the audio and the text - I will do the rest.

Models: Qwen3-ASR-1.7B: ready, Qwen3-ForcedAligner-0.6B: ready, Qwen3-TTS-12Hz-1.7B-Base: ready, sage-fred15-distilled-95m: ready

Use voices responsibly: only with the voice owner's permission and in line with their rights. Recordings and results stay on this computer.

The Train window, lower part: quick preview, name and description, gender and age, the universal model, the dataset button and the progress stages.

When the job ends, the status shows **The voice was added to "My voices"**. and a button **Open the result folder** appears. The training folder is `Projects\Voices\<voice name>_Voxprint` in the projects folder: the dataset is in `dataset`, the adapter (tens of MB) in `output\<voice name>`.

The pictures show the default mode with a text: **Choose audio** and **Choose text** one under the other, each with the name of the chosen file. Ticking **I have no transcript - recognise the speech automatically** switches to the audio-only mode (section 6.4): the text row becomes an optional script and a warning block appears.

6.3 Training presets

Training quality offers four presets. For each one the program shows what it means and an **estimated time** – for example "about 2 min on NVIDIA GeForce RTX 4090 (5 passes, adapter size 32, about 71 clips)". Read the estimate as "about".

Preset	Meaning
Fast	A quick try: fewer passes over your recording and a smaller adapter. The voice is recognisable, fine details are rougher.
Balanced (recommended)	The automatic choice for your recording and graphics card. Good likeness in reasonable time.
Maximum	The longest run: about twice the passes and a bigger adapter. Best chance of a close voice, but it takes longest and may overfit on short recordings.
Manual	You choose the numbers yourself (see Advanced). Only change them if you know what they do.



Under **Advanced** you see the numbers behind a preset:

Field	Plain meaning
Passes over the recording (epochs)	How many times the training goes over your whole recording. More passes: closer to the voice, but too many can make it "babble".
Adapter size (LoRA rank)	How big the voice add-on is. Larger – more detail, more memory.
Adapter strength (LoRA alpha)	How strongly the add-on influences the base model.
Learning rate	How big each learning step is. The most sensitive number: do not raise it unless you know why.
Steps per update (gradient accumulation)	How many small steps are summed before the model is updated (saves memory).

Editable only in Manual. Higher learning rates and many passes can lower the training number yet make the voice babble and never stop - the automatic value is safest. No preset ever raises the learning rate. In our tests a lower "training number" (loss) did not mean a better voice – always listen to the result.

6.4 No transcript? (audio only)

Tick **I have no transcript - recognise the speech automatically** and the program recognises the speech itself (Qwen3-ASR, offline; the recognition model, about 1.9 GB, is part of the first-start download). The text row stays as an **optional script** (**Script (optional): the text that was read aloud...**, or drop a **.txt** into the window): if you choose the text that was read aloud, every recognised piece is matched to it tolerantly (stumbles and re-read lines are handled), so the final text is the correct one, not the recognised one. In addition you have:

- **Choose audio** – now you can pick **many files** at once; the line shows "N audio file(s) chosen";
- **Choose a folder with clips...** – a whole folder of clips, searched with its subfolders;
- drag and drop of files into the window.

Each file is recognised separately; long recordings are cut at pauses into pieces of up to 14 s; everything is merged into one dataset and made equally loud. Skipped automatically: clips under 1.5 s or over 20 s, empty or implausible recognition, duplicates and clips with poor sound. At the end the window reports, for example, "Recognised N file(s): kept X of Y clips = ... min of ... min of audio".

Automatic speech recognition can make mistakes, and the result is not guaranteed: wrong words in the text teach the voice wrong sounds. Clips the program is unsure about are skipped, but please use clear recordings of one speaker. You can pick many files (or a folder) at once, or drop them here.

Because of this, you must tick **I understand that the text will be recognised automatically and may contain errors** before the buttons work (the tick resets every time you switch the mode on).

6.5 Quick preview and comparing two variants

Before a long run you can test settings cheaply. **Quick preview...** trains a very short adapter on a few clips (at most 12 clips and 4 passes, about 3 minutes at most) and then makes a sample of about 10 seconds that you can listen to in the player. The estimate is shown next to the button: **Quick preview: about ... (... short training run(s) on a few clips, then a ~10 s sample with automatic checks).**



The small adapter (tens of MB) works only in a few apps, e.g. Alexandria.

☒ Compare 2 variants

☒ Check the voice after training (recommended)

Quick preview: about 2 min (2 short training run(s) on a few clips, then a ~10 s sample with automatic checks).

Listen

Variant A: 4 passes, adapter size 32 - sample 22 s, recognition errors 15 %, pitch vs. your recording -0.7 semitones - looks good

Variant B: 6 passes, adapter size 64 - sample 32 s, recognition errors 15 %, pitch vs. your recording +1.2 semitones - looks good

Quick preview of two variants, with the player and the automatic metrics.

- With **Compare 2 variants** ticked, two samples are made: A with the current settings and B with a larger adapter and 1.5 times the passes.
- Every variant gets a line of automatic metrics: the length of the sample, "recognition errors" (how many words an automatic listener got wrong when re-reading the sample – lower is better) and "pitch vs. your recording" in semitones (how far the voice is higher or lower than yours – close to 0 is best). A verdict follows: **looks good**, or a warning.
- **Play** plays a variant. **Use these settings** copies its numbers into the **Manual** preset; then press **Create voice (LoRA)** for the full run.
- A preview voice is temporary and is **not** added to **My voices**. The metrics are a hint, not a verdict: always listen.

6.6 The consent card

Every voice carries a record of what its owner allows. Under **Voice owner's consent** choose how it is set:

Option	When to use it
Read it from the spoken statement at the end of my recording (recommended)	You finished the recording with one of the three consent sentences (block 18 of the script). After training the program recognises the statement and shows the result.
I will choose the scope myself	You type the speaker's name (placeholder " Speaker's name ") and choose the scope yourself. Use it only for your own voice or with the owner's permission.
No statement: private use only	No statement. The voice is marked <i>private only</i> .

The check box **Keep the spoken statement with the voice** stores the spoken statement (the last 45 seconds of the recording) next to the voice.

After training in the automatic mode the **consent card** appears:



☒ Keep the spoken statement with the voice

Detected consent - name: Aleksandr Mitroshenkov; date: 2026-10-03; scope: Public, non-commercial: may be published free, not sold.

Recognised statement: "I, Aleksandr Mitroshenkov, today, October third, twenty twenty-six, allow this recording and the characteristics of my voice to be used to train neural networks and to synthesize speech, including free publishing but not selling."

Check it and confirm (or change the scope) with one click.

Public, non-commercial: may be published free, not sold ▼ Confirm

The consent card after training: detected name, date and scope, and the Confirm button.

It shows the *detected* name, date and scope, and the sentence that was recognised. Check it, change the scope in the drop-down if needed and press **Confirm**. If the statement is unclear or missing, the strictest level – **Private only** – is used. This is a record of what was said, **not legal advice**.

Scope in the drop-down	Badge	Meaning	Licence stored
Commercial: the audio may be sold	Commercial OK	The voice owner allowed commercial use: narration made with this voice may be sold.	CC-BY-4.0
Public, non-commercial: may be published free, not sold	Public, non-commercial	Narration made with this voice may be published free of charge, but not sold.	CC-BY-NC-4.0
Private only: results stay on this computer	Private only	Narration made with this voice must stay on your computer; do not publish or sell it.	custom/ personal-only

6.7 The universal model (optional)

The small adapter (tens of MB) works only in a few apps (for example Alexandria). **Build universal model (~4 GB)** merges it into a standalone model folder (`merged_model`) that works in **any** app that runs Qwen3-TTS. Voxprint checks the free disk space and asks first (**Build the universal model**); it is never built automatically. The button is disabled until a voice has been created. Details are in `USAGE.txt` inside the folder.

7 My voices

Open this from the **My voices** card. **Voices are stored on this computer. Each voice has a licence that says what its narration may be used for.**



← Studio
My voices

Voices are stored on this computer. Each voice has a licence that says what its narration may be used for.

Import voice...
Download voices from repository
Refresh list

Boaz
CC0-1.0 · commercial use OK
Commercial OK

Russian · 28 min of speech · 2 epochs · Male voice · by Mark Chulsky (LibriVox) · public domain

Bundled open male Russian voice, trained only from the public-domain LibriVox recording 'Аграфена (Agrafena)' by Boris Zaytsev, read by Mark Chulsky (<https://librivox.org/agrafena-by-boris-zaytsev/>). Free for any use, also commercial (CC0-1.0). Attribution appreciated. The name Boaz is Voxprint's own and does not imply the reader's endorsement.

Comes with Voxprint (read-only). Licence: CC0-1.0 - free for any use, also commercial.

Preview
Narrate with this voice
Details...
Export...
Open folder
Delete

Tirzah
CC0-1.0 · commercial use OK
Commercial OK

Russian · 29 min of speech · 2 epochs · Female voice · by Anastasiia Solokha (LibriVox) · public domain

Bundled open female Russian voice, trained only from the public-domain LibriVox recording 'Степные сказки (Stepnyia skazki)' by Grigory Danilevsky, read by Anastasiia Solokha (<https://librivox.org/stepnyia-skazki-by-grigory-danilevsky/>). Free for any use, also commercial (CC0-1.0). Attribution appreciated. The name Tirzah is Voxprint's own and does not imply the reader's endorsement.

Comes with Voxprint (read-only). Licence: CC0-1.0 - free for any use, also commercial.

Preview
Narrate with this voice
Details...
Export...
Open folder
Delete

Open universal voice
CC0-1.0 · commercial use OK
Commercial OK

English · by Linda Johnson (LibriVox) / Keith Ito (LJ Speech) · public domain · 54 MB

Open universal voice: trained only from public-domain speech (LJ Speech 1.1, reader Linda Johnson, LibriVox). Free for any use, also commercial. Attribution appreciated: LJ Speech dataset, Keith Ito, <https://keithito.com/LJ-Speech-Dataset/>

Not downloaded yet

Download

The licence covers the voice model only. Use a voice only with its owner's permission and within its licence (for example CC BY-NC = personal use only; you are responsible for lawful use.

My voices: the bundled voices Boaz and Tirzah (CC0-1.0, read-only) and the Open universal voice, which is not downloaded yet.

Each voice is a card:

Element	Meaning
Name and two badges	The licence badge (green "... · commercial use OK" or amber "Personal use only · ...") and the scope badge (Commercial OK , Public, non-commercial , Private only).
Information line	Language · minutes of speech used for training · epochs (passes) · voice type · author · (for online voices) size.
Description	What the author wrote. A "test use only" voice also shows a warning.
Preview	Plays a short sample of the voice. Stop stops it. If the voice has no sample: <i>This voice has no sample</i>
Narrate with this voice	Opens the Narrate window with this voice already chosen.
Details...	Opens Voice details : Name , Author , Licence of the voice , Gender and age , Description . Press Save .
Export...	Saves the voice for another computer: Voice file (~... MB) (the trained voice; another Voxprint imports it in a second) or Full model (~4 GB) (the universal model).
Open folder	Opens the voice's folder.
Delete	Deletes the voice from this computer after a confirmation: <i>Delete the voice "..."</i> <i>from this computer? This cannot be undone.</i> Bundled voices are read-only: Details... and Delete are greyed out (<i>A voice that comes with Voxprint cannot be edited or deleted.</i>).

The footer of the window repeats: ***The licence covers the voice model only. Use a voice only with its owner's permission and within its licence (for example CC BY-NC = personal use only); you are responsible for lawful use.***



When the library is empty the window says **No voices yet. Train your own voice from a recording, import a voice folder or .zip, or download voices from the repository.** and offers **Train your voice**.

7.1 Online voices and the voice repository

The third card in the picture above, the *Open universal voice*, is such an online voice: **Not downloaded yet** and a **Download** button.

- Voices from the online index appear by themselves as cards with the state **Not downloaded yet**, their licence and size and a **Download** button. The download is checked by a SHA-256 checksum and can be resumed.
- **Refresh list** reloads the list. Offline, Voxprint shows the last known list (**No connection - showing the last known list of online voices.**).
- **Download voices from repository** opens the **Voice repository** dialog with **Refresh**, a field for the address of the index (`index.json`) and **Download selected**.

7.2 Importing a voice

Import voice... offers two ways: **From a folder...** (an adapter folder) and **From a .zip file...** (a voice package). Archives are checked for unsafe paths and size limits. A voice that arrives without a declared licence is treated as *personal use only*.

7.3 Licences and the voice that Voxprint offers

Licence	Commercial use
CC0-1.0, CC-BY-4.0, CC-BY-SA-4.0	allowed (keep attribution / share-alike if the licence asks for it)
CC-BY-NC-4.0, CC-BY-NC-SA-4.0	not allowed
custom/personal-only (default)	not allowed; results stay on your computer
custom/test-use-only	not allowed; for testing only, no public release

- **Boaz** (male) and **Tirzah** (female) – two open **Russian** voices that come with Voxprint (part of the standard model download, about half an hour of speech each). They are trained only from public-domain LibriVox recordings: Boaz from *Аграфена (Agrafena)* by Boris Zaytsev, read by Mark Chulsky; Tirzah from *Степные сказки (Stepnyia skazki)* by Grigory Danilevsky, read by Anastasiia Solokha. Licence **CC0-1.0** – free for any use, also commercial; attribution is appreciated. The names are Voxprint's own and do not imply the readers' endorsement. In the library they are read-only (**Comes with Voxprint (read-only). Licence: ... - free for any use, also commercial.**). Samples – Genesis 1:1–2:3 (Russian Synodal translation, public domain) narrated with both voices, also CC0 – are in the repository folder `samples/`.
- **Open universal voice** (German UI: *Offene Universalstimme*, Russian UI: *Открытый универсальный голос*) – an English voice for people who have no recording of their own. Trained only from the public-domain LJ Speech dataset (reader Linda Johnson, LibriVox); licence **CC0-1.0** – free for any use, also commercial; attribution to the LJ Speech dataset (Keith Ito) is appreciated.


The Open universal voice is downloaded separately from the voice repository.

The pictures in this manual show the bundled voices. Voices you train yourself are *personal use only* by default (scope **Private only**): no public use of the output, no commercial projects.

The badge is information, not legal advice. The licence covers the voice model only.

8 Settings

Open **Settings** with the gear button (top right of the Studio and the other windows).



Settings (build 667; the projects path and the network address are blurred).

Control	What it does
Language	The interface language: English, Deutsch, Русский, Українська, Latviešu. It changes at once in all windows.
Window transparency	How translucent the windows are: Default , More or Off .
Network interface	Which network adapter the downloads use. Automatic : if a download cannot connect (VPN, unusual adapter), the other interfaces are tried and the working one is remembered.
Speech recognition model	Which speech-recognition model is used. Automatic (by video memory) takes the 1.7B model on cards with about 8 GB of video memory or more, otherwise 0.6B. The line below the check box shows the model in use.
Preload models into memory at startup	After the start, loads the installed speech models into RAM in the background, so that narration starts at once. The available RAM is shown next to it.
Narration: pauses and speed	Pause after a comma , Strong break (dash, colon, semicolon, \", and / but\") , Pause after a sentence , Paragraph or verse line , Chapter, title or scene break – the five pause lengths in seconds (defaults 0.25 / 0.40 / 0.60 / 1.00 / 2.00 s); Reading speed (70–130 %; 100 % is the voice's own speed); Reading style . See section 5.4.



Control	What it does
Default pauses and speed	Restores the default pauses and speed.
Check for updates	Checks whether newer verified versions of components or models exist and installs them (with a safety check and automatic rollback). Voxprint also checks quietly once a week. Messages: <i>Everything is already up to date.</i> , <i>Restart the program to start using the updates.</i>
Open the models folder	Opens the folder with the downloaded models.
Open the data and log folder	Opens the data and log folder.
Save diagnostic report...	Saves the log files, information about this computer (system, GPU, memory) and the settings as one zip file; the names of your files are cut out. Attach it when you report an error.
Components	Opens the Components window: what is installed, updates, and the complete download.
Projects folder	Where new projects (audiobooks, voice trainings, Re-voice) are kept. Open projects folder , Change... (prefer a local drive with plenty of space that is not synced to the cloud), Default . Existing projects stay where they are.
Models and voices – Include models , Include my voices	What goes into a backup (both on by default).
Use models from this folder (don't copy)	When restoring, read the models from the backup drive instead of copying them. The drive must stay connected.
Back up models & voices...	Copies models (and voices) to a folder or drive you choose. Shows the size and free space first; resumable; identical files are skipped; every copy is checked. A progress bar and a Cancel button appear.
Restore from backup... , Restore from backup folder	Restores from a Voxprint-backup folder and checks every file by its checksum. Voices that already exist with different content are never overwritten; damaged or missing files are downloaded as usual.
Existing models folder	Models from a previous install or a backup: they are linked (same drive) or copied into Voxprint and checked by checksum before anything is downloaded. Choose folder... picks the folder, Clear forgets it.
Check and repair – Check & repair	Checks the program, its components and all models file by file (SHA-256 checksums) and downloads whatever is missing or damaged from the internet again (only the bad files). Your voices and books are not touched. See section 8.2.
About	Version, the idea, how it works, open-source components with their licences, a link to the GitHub repository and third-party notices.

The bottom line shows the version, for example *Voxprint AI Audiobook Builder 0.1.3-beta · build 667 "Menuchah"*.

If the program notices a damaged installation, a banner offers the repair (**Voxprint's installation needs repair: ...**). If **Check & repair** does not help, run the Voxprint installer again – voices and models are kept.



When the program finds an older component in an environment it does not own (for example your own Python), it asks first: **Newer component available** with **Upgrade** or **Not now**; it never changes your environment silently.

8.1 Maximum quality out of the box


Out of the box, everything that can run automatically is **already selected** and marked with a star and the word **recommended**: if you touch nothing, Voxprint works at maximum quality. The models these options need are part of the first-start download, so there is nothing extra to click. The single check boxes stay in their windows.

- Text preparation (Narrate a book): the seven rule-based steps – layout, footnote marks and page numbers, quotes and dashes, links, chapter headings, numbers in words, abbreviations.
- AI clean-up of typos (Russian books; one-time download of about 365 MB).
- Translation of the book before narrating (English, Russian, German; Opus-MT models, about 300 MB per direction, downloaded once).
- Training window: **Check the voice after training (recommended)** and **Compare 2 variants** (both need the speech-recognition model, about 1.9 GB, downloaded once).
- Always on, without a switch: alignment of the text onto the audio and its plausibility check, the audio quality filter, the recognition filters of the audio-only mode.
- Not included, because they do not exist yet: punctuation model, stress marks, speaker roles.

The training preset stays at **Balanced (recommended)** on purpose: in our tests longer training with a higher learning rate made voices babble, so more is not better there. Instead, after **Compare 2 variants** the better of the two variants is marked **recommended** (verdict first, then fewer recognition errors, then the smaller pitch shift; with no clear difference variant A, which is cheaper). The choice is still yours: always listen.

8.2 Check and repair


Check & repair checks the program, its components and every model file by file against its SHA-256 checksum and downloads only what is missing or damaged. Your voices and books are not touched. While it runs, a progress bar and a step line (for example *Step 8 of 21: Checking Qwen3-ASR-0.6B: model.safetensors*) show where it is, and the button becomes **Stop the check**; a stopped check keeps what it has finished.



Check and repair while it runs: the progress bar and the current step.



At the end a summary appears, for example *Check finished: 24 checked, 0 repaired or downloaded, 0 failed*. Whatever could not be repaired is named below it.



Check and repair finished: the summary line.

Known issue in build 667: while the check runs, the right column of Settings is squeezed and the explanatory texts overlap (first picture). This is cosmetic and will be fixed in a coming build.

The same check runs without the window as `Voxprint.exe --auto-repair` (chapter 11).

9 Glossary

Adapter (LoRA)

A small add-on file (tens of MB) that "teaches" the big speech model one particular voice. Technically LoRA – low-rank adaptation. The button **Create voice (LoRA)** creates it.

Universal model

The adapter merged into a full ~4 GB model folder. It works in any app that runs Qwen3-TTS.

Qwen3-TTS

The open speech-synthesis model that Voxprint uses to speak. The base model is Apache-2.0 licensed.

Dataset

The training material Voxprint cuts from your recording: short clips (3–12 s) with their texts, a reference clip and a report.

Alignment



Finding out at what second each word of the text is spoken in the recording. It allows clean cutting at pauses, never in the middle of a word.

Clip / fragment

A short piece of the recording in the dataset. In the narration window "Fragment 3 of 8" means a piece of the book.

Chunk

A piece of the book text of about a sentence or a clause that is sent to the voice model in one go. The UI calls it a "fragment". Finished chunks are saved, which makes narration resumable.

Batch

Several chunks that the voice model makes together in one call. Bigger batches are much faster but use more video memory (up to 12 at a time).

RTF (real-time factor)

Seconds of computing per second of produced audio. 0.31 means one hour of audiobook takes about 19 minutes. Lower is faster.

Epoch / pass

One complete pass of the training over your whole recording (UI: " **Passes over the recording (epochs)** "). The voice list shows how many epochs a voice was trained for.

Preset

A ready-made set of training numbers: **Fast**, **Balanced (recommended)**, **Maximum**, **Manual**.

Rank, alpha, learning rate, gradient accumulation

Technical training numbers shown under **Advanced**; see section 6.3.

Preview

In **My voices**: a short sample of a voice (**Preview**). In the Train window: the *quick preview* – a short test training with a ~10 s sample and automatic checks; nothing is added to your voices.

Recognition errors (WER)

The share of words that an automatic listener (Qwen3-ASR) gets wrong when it transcribes the generated sample. Lower is better.

Pitch, semitone

How high the voice is. A semitone is the smallest step on a piano. "Pitch vs. your recording: +1.2 semitones" means the voice is slightly higher than the original.

Consent

The voice owner's permission, spoken at the end of the recording or entered by hand, saved with the voice. A record, not legal advice.

Scope

What the consent allows: **Commercial OK**, **Public, non-commercial** or **Private only**.

Licence badge / scope badge

Coloured labels on a voice card and in the Narrate window. Green: commercial use allowed. Amber: restricted. The badge is information, not legal advice.

Voice type

An optional label of the voice: male, female, child or other. Saved with the voice.

Mini player

The small player in the Narrate window that plays the finished part of the book while the rest is still being made.

**Bitrate (kbit/s)**

How much data one second of sound takes. Higher means better quality and bigger files. Speech needs little: Opus 32 kbit/s is plenty.

Opus, MP3, AAC / M4B, FLAC, WAV

Audio formats. Opus: small, open, free. MP3: plays everywhere. AAC in an M4B container: for Apple Books (patent-encumbered). FLAC and WAV: lossless, large files for archiving and editing.

Chapter marks

Bookmarks inside a single audio file that let a player jump from chapter to chapter.

Video memory (VRAM)

The memory of the graphics card. It limits the size of the model and of the batch.

Resume

Continuing a stopped narration or download from where it stopped.

Checksum (SHA-256)

A fingerprint of a file that proves it was downloaded or copied without damage.

Mirror (ModelScope)

An alternative download server that Voxprint uses if Hugging Face is slow or unreachable.

Check and repair

Checks every component and model file by its checksum and downloads only what is missing or damaged (**Settings** → **Check & repair**, or `--auto-repair`).

Measured pauses

Silence of a fixed length that Voxprint inserts between the spoken pieces – after a comma, a sentence, a paragraph, a chapter – independent of how the voice model phrases the text.

Bundled voices

Boaz and Tirzah: two open Russian voices (CC0-1.0) that come with Voxprint and are read-only in the library.

10 Troubleshooting and FAQ

Which messages can I get, and what do I do?

Message	What to do
The recording is too short (under 10 seconds). At least a few minutes are needed.	Record longer – 5–15 minutes.
Could not read the audio file. Try another format (WAV or MP3).	Convert the recording to WAV or MP3.
The recording and the text do not seem to match: the text is much longer or shorter than the recording. Check that you chose the right files and read the whole text.	Check that you chose the right audio and text and read the whole text.
The recording has long stretches of speech that are not in the text (at the beginning or the end). Trim the extra part or choose another text.	Trim the extra speech from the start or end of the recording, or choose the correct text.



Message	What to do
All fragments of the recording turned out too quiet or distorted. Check the microphone and record again.	Check the microphone and level; record again in a quieter place.
Not enough video memory (VRAM).	Close other programs that use the graphics card; Voxprint retries with a lighter load (<i>Not enough video memory - trying again with a lighter load...</i>). As a last resort Continue on the processor (slow) .
Could not download the model {short}. Check your internet connection and try again. You can download the model manually: {url}	Check your internet connection and start again; the download continues where it stopped.
The voice engine failed on fragment {n}. Finished fragments are saved - start again to continue.	Press Start narration again – finished fragments are kept.
ffmpeg is needed for this format but was not found. Use “Repair” in Settings or choose WAV.	Settings → Check & repair , or choose WAV.
Unsupported book format. Use TXT, FB2 (or FB2.ZIP) or EPUB.	Convert the book to TXT, FB2 or EPUB.
The book file looks unsafe or is too large and was not opened.	The file looks damaged or too large; try another copy.
No NVIDIA graphics card is being used: training will run on the processor and take many hours. A NVIDIA graphics card is needed for a fast result.	Training works without an NVIDIA card but is extremely slow.
The final training loss is low ($\{loss\} < \{threshold\}$): the voice may have turned to “mush” or the speech may not stop. If it sounds bad, record more text and try again.	The voice may babble. Use Balanced (recommended) , record more text, and listen with the quick preview.

The voice sounds higher than mine. A small upward pitch drift (about +2...+3 semitones was seen on a male voice) is a known effect of this kind of training. The automatic check warns above +4. Try the quick preview with both variants and take the one that is closer.

The voice babbles or never stops. Too many passes or too high a learning rate. Use **Balanced (recommended)** , never raise the learning rate by hand.

Can I stop and continue later? Narration: yes, always (section 5.5). Downloads and backups: yes. Training: cancel is possible, but a full run starts over; use the quick preview first.

Can I use my voice for a published audiobook? Only if the consent scope allows it (**Commercial OK** to sell; **Public, non-commercial** to publish for free). Voices you train without a consent statement are *private only*. You can change the licence under **Details...** if you own the voice.

Does Voxprint send my recordings anywhere? No. It sends nothing to the internet except downloading models, checking for updates and, when you open the repository, fetching the voice index and the voices you choose. No telemetry, no account.

Which languages are supported? The interface: English, Deutsch, Русский, Українська, Latviešu. Text preparation covers Russian and English fully. The speech model supports further languages, but they are not tested.

Where is the log? **Open the data and log folder** in Settings. Send the file from there if an error repeats.



Is the AAC / M4B option safe to use? It is optional and off by default for good reason: AAC is patent-encumbered and you are responsible for compliance (section 5.4).

11 Command line

Everything in this chapter is optional: the windows do all of it. The command line is meant for scripts, servers and AI agents and uses the same pipeline as the windows. The full reference is `docs/CLI.md`; the file to hand to an agent is `docs/AGENTS.md` (both in the repository).

11.1 Where the program is and the rules

The installed program is `C:\Program Files\Voxprint\Voxprint.exe` (another folder if you chose one; `where.exe Voxprint` finds it). In a source checkout use `python main.py`. Below, `voxprint` stands for either:

- Nothing prompts; `--yes` is accepted by every command.
- `--json` prints progress as JSON lines and ends with one result object (`type: result`, `ok`, `exit_code`, `outputs`, `warnings`, `error`, `hint`).
- `voxprint --help` and `voxprint <command> --help` show copy-paste examples; `voxprint --version` prints, for example, `Voxprint 0.1.3-beta build 667 "Menuchah"`.
- Set `VOXPRINT_LANG=en` if you parse the messages of the pipeline.

Exit code	Meaning
0	Success
1	Unexpected failure – save <code>voxprint diag</code>
2	Missing or unknown arguments
3	Book, audio, transcript, voice or backup missing or unreadable
4	A model or component is not installed, or its download failed
5	Out of GPU memory – for training, retry with <code>--force-cpu</code>
6	Cancelled – run the same command again to resume

11.2 Narrate a book

```
voxprint narrate BOOK --voice ID_OR_NAME --out DIR [--format mp3,m4b,opus,...] [--pause-comma S] [--pause-mid S] [--pause-sentence S] [--pause-paragraph S] [--pause-chapter S] [--speed X] [--style auto|scripture|fiction|dialogue] [--pauses] [--ai-disclosure] [--json]
```

The book is a TXT, FB2, `.fb2.zip` or EPUB file; the result lands in `<DIR>/<book title>/`; the default format is one Opus file. The pause flags (seconds), `--speed` (0.7–1.3) and `--style` override the values from **Settings** → **Narration: pauses and speed** for one run; `--pauses` is the opt-in cut at every comma. Running the same command again skips the fragments that are already cached, and changing pauses or speed never synthesizes them again.

```
voxprint narrate genesis.txt --voice Boaz --out ./audiobooks --format mp3 --json
```



11.3 Train a voice

```
voxprint train AUDIO [--text SCRIPT] [--name NAME] [--type male|female|child|other] [--out DIR]
[--consent none|auto|commercial|public_noncommercial|private_only] [--speaker NAME] [--license
ID] [--language CODE] [--force-cpu] [--json]
```

With `--text` the recording is aligned to the script; without it (one file or a folder of clips) speech recognition builds the dataset. `--out` is the parent folder; every voice gets its own `<voice name>_Voxprint`. Without `--consent` the voice is stored as *private only*; `--consent auto` reads the spoken statement, the other values record a scope you confirmed yourself, and `--license` can never allow more than the consent. Only train a voice you have the right to use.

```
voxprint train recording.wav --text script.txt --name Anna --type female --consent auto
```

11.4 Back up and restore

```
voxprint backup --out DIR [--no-models] [--no-voices] [--json] voxprint restore --from DIR [--
link] [--json]
```

`backup` copies the models and the voice library into `<DIR>/Voxprint-backup/` with the manifest `voxprint-backup.json` (version, build, size and SHA-256 of every file). Files already there with the same size and hash are skipped, so an interrupted copy continues; nothing is written when the drive is too small. `restore` copies the backup into the normal folders and checks every file; a damaged or missing file is named and later downloaded as usual. `--link` keeps the models on the backup drive (it must stay connected) and copies only the voices. Errors exit with code 3. These are the same jobs as the buttons in **Settings** → **Models and voices**.

11.5 Check the installation

```
voxprint status --json # version, GPU, voices, formats, installed models voxprint models
download required # TTS, aligner and speech recognition, if missing Voxprint.exe --verify-
install # quick install check with reason codes Voxprint.exe --auto-repair # Check & repair:
every file by SHA-256 voxprint diag --out report.zip # diagnostic report
```

`status` (also `capabilities`) prints one JSON object and downloads nothing – run it before `train` or `narrate`. `--auto-repair` is the same job as **Settings** → **Check & repair** (section 8.2): it prints one line per checked item, downloads only the damaged or missing files and exits with 0 when everything is in order; the report is also written to `logs\auto_repair.txt`. `--verify-install` writes `logs\verify_install.txt`. `diag` is the same as **Save diagnostic report...**.

12 Licences and responsibility

- **Voxprint's own source code is licensed under the Apache License 2.0** (Copyright 2026 Aleksandr Mitroshenkov; see the `LICENSE` and `NOTICE` files).
- **The Apache licence of the code does not cover models and voices.** The models Voxprint downloads (Qwen3-TTS, Qwen3-ASR and the optional text models) keep their authors' licences. Every voice has its own licence, shown as a badge and stored in `voice.json`.
- Third-party components, with their licences, are listed under **About** and in `THIRD_PARTY_NOTICES.md` and the `licenses` folder of the installation. Qt/PySide6 is used under the LGPL-3.0; the bundled ffmpeg is a separate program with its own licence; the AAC codec is patent-encumbered.



- **Use voices responsibly.** A voice is personal data and a personality right. Use a voice only with the explicit permission of its owner and within its licence; no impersonation, fraud or misleading content; label synthetic speech where this is required. You are responsible for the lawful use of the voices and audio you create.

This manual and the licence badges are information, not legal advice. The consent card records what a speaker said; it does not replace the law of your country.

Voxprint AI Audiobook Builder 0.1.3 (beta), build 667 "Menuchah". Manual of 9 October 2026.

Voxprint AI Audiobook Builder is an independent project and is not affiliated with other products or services with similar names.